

Fast boot and suspend/resume

Zhang Rui
rui.zhang@intel.com

Beijing October 18th 2008



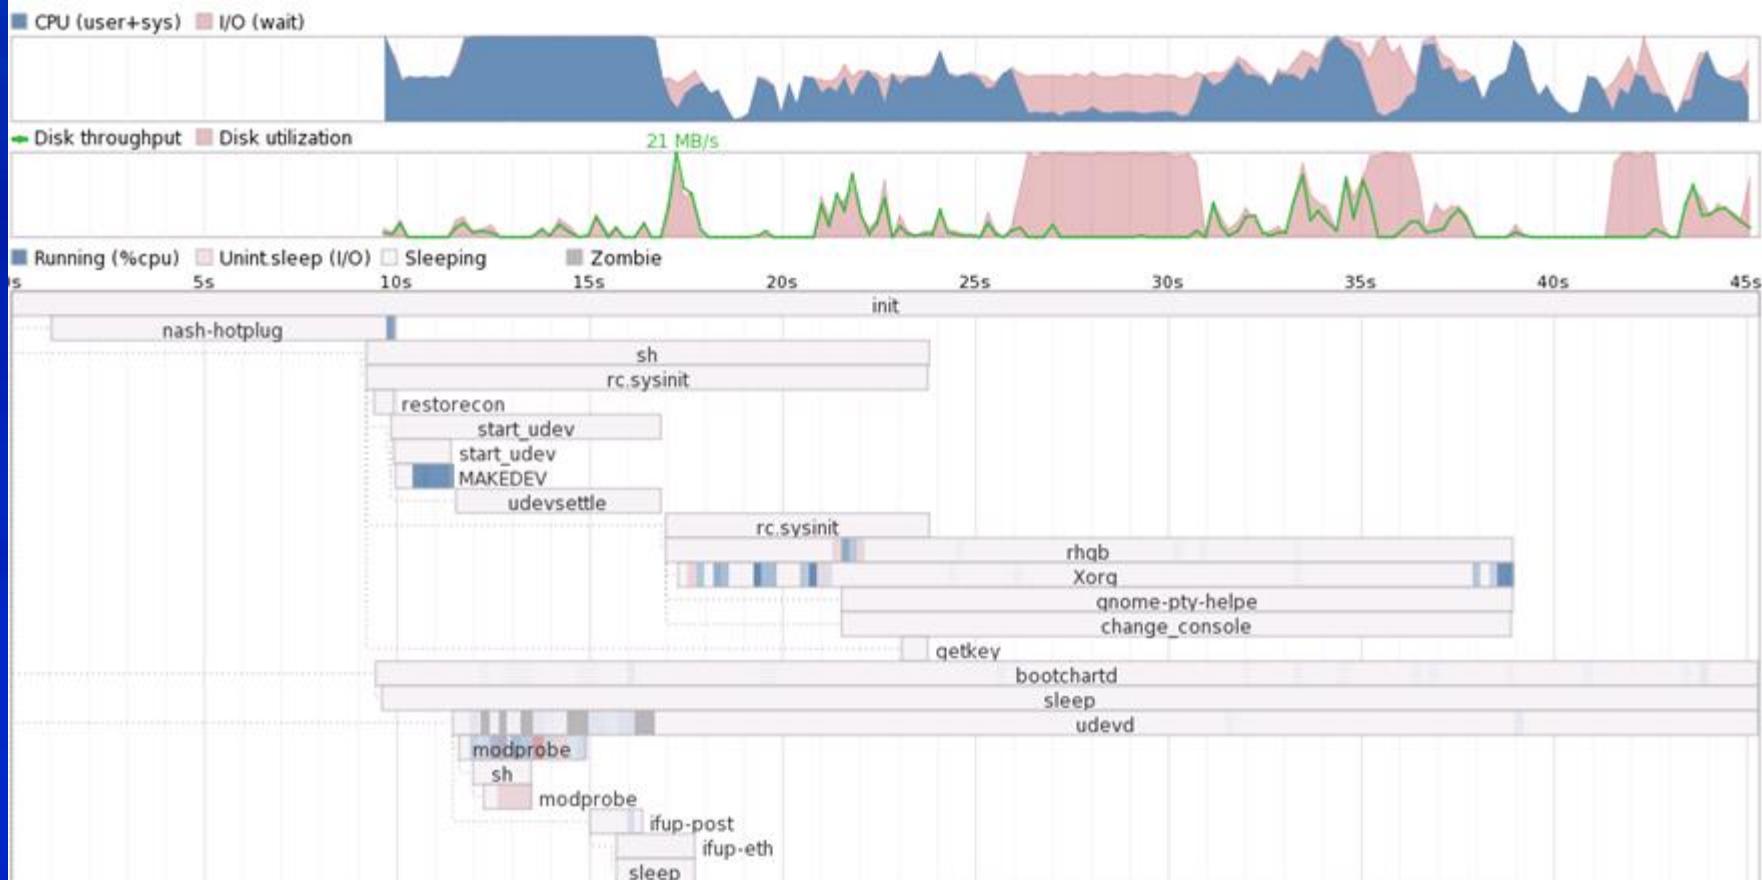
Agenda

- **Fast boot**
 - Current State of the Art
 - Splash Screens
 - Budget & Resource Scheduling
 - Kernel
 - Early Userspace
 - X.org
- **Fast suspend/resume**
 - Driver optimization
 - Device asynchronous resume
 - More aggressive attempts

Current state

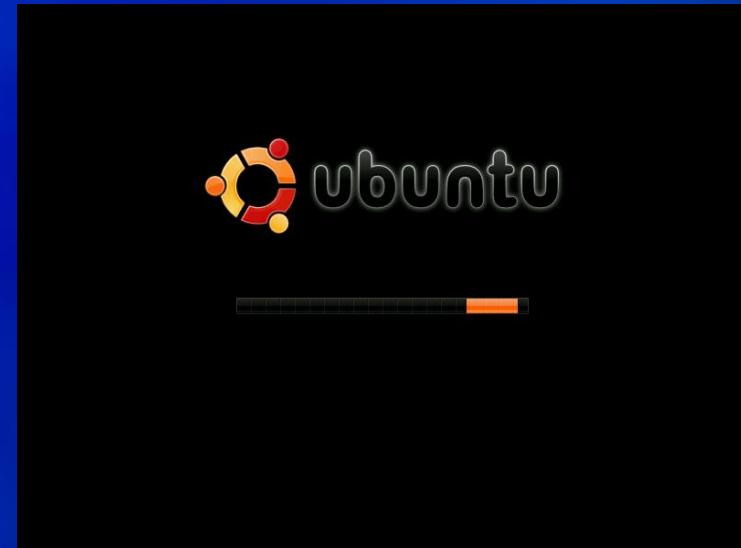
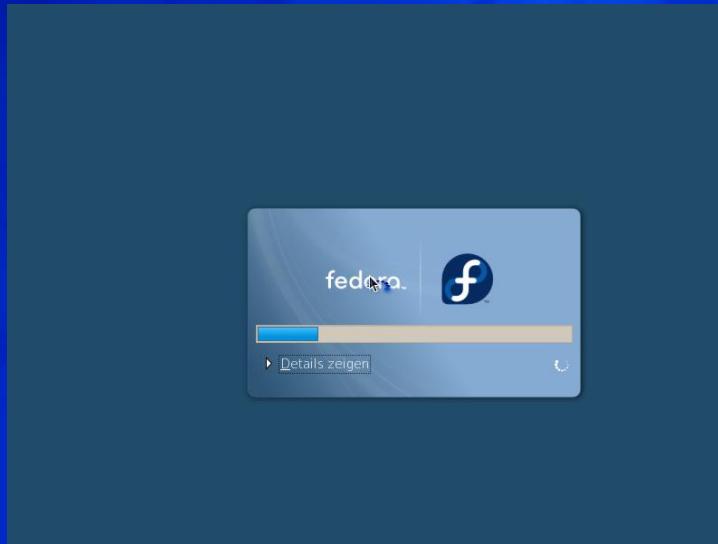
Boot chart for localhost.localdomain (Tue Sep 2 05:35:37 PDT 2008)

uname: Linux 2.6.25.14-108.fc9.i686 #1 SMP Mon Aug 4 14:08:11 EDT 2008 i686
release: Fedora release 9 (Sulphur)
CPU: Intel(R) Atom(TM) CPU N270 @ 1.60GHz (1)
kernel options: ro root=UUID=140d85e0-58e5-4a20-a3ce-c325402c326e rhgb quiet init=/sbin/bootchartd
time: 0:45



We hate splash screens

By the time you see it.... we want to be done!



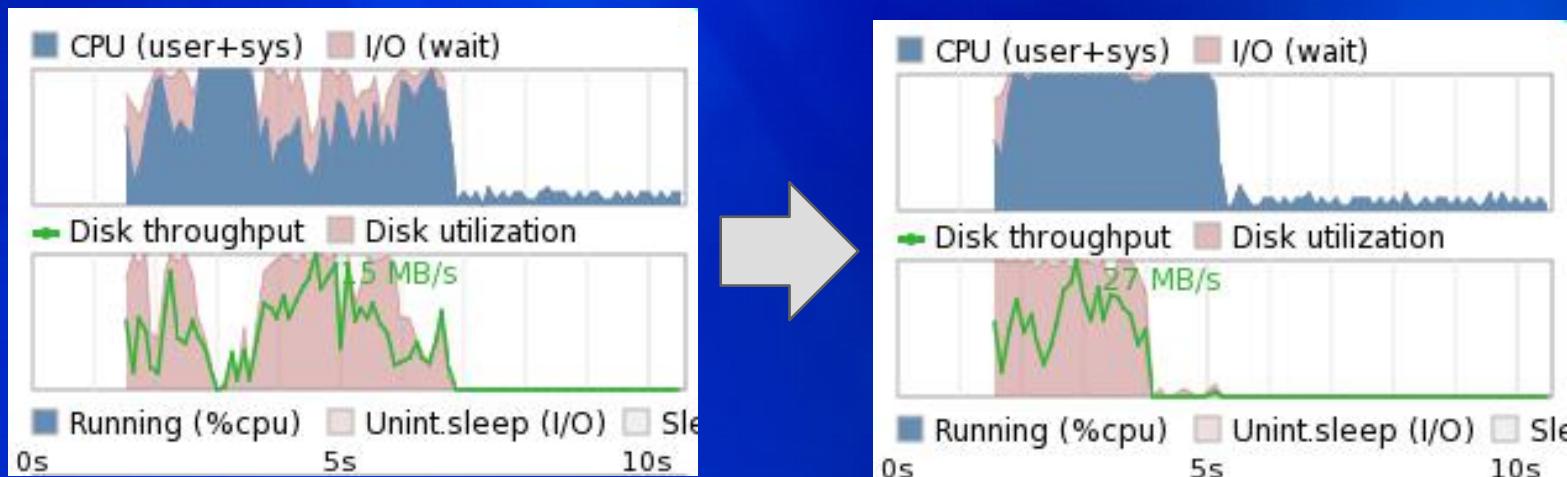
Time budgets, Resource scheduling

boot to be completely done in 5 seconds

Component	Allocated time
Kernel	1 second
Early boot	1 second
X	1 second
GUI/Desktop	2 seconds

sReadAhead

- read used portions of files in "use order" as early as possible to prime the pagecache
- "idle" IO scheduler class



Kernel

Budget: 1 second

- All system components built into the kernel image
 - Modules are slow, synchronous, and not needed for core components
- No initrd
 - All key drivers are in the kernel
 - /dev is populated with the fixed device nodes
 - initrd management just plain takes too long



Kernel

Budget: 1 second

- **Asynchronous initialization of non-essential components with a new, asynchronous initcall level**
- **Reference:**
<git://git.kernel.org/pub/scm/linux/kernel/git/arjan/linux-2.6-fastboot.git>



Early boot / Init

Budget: 1 second

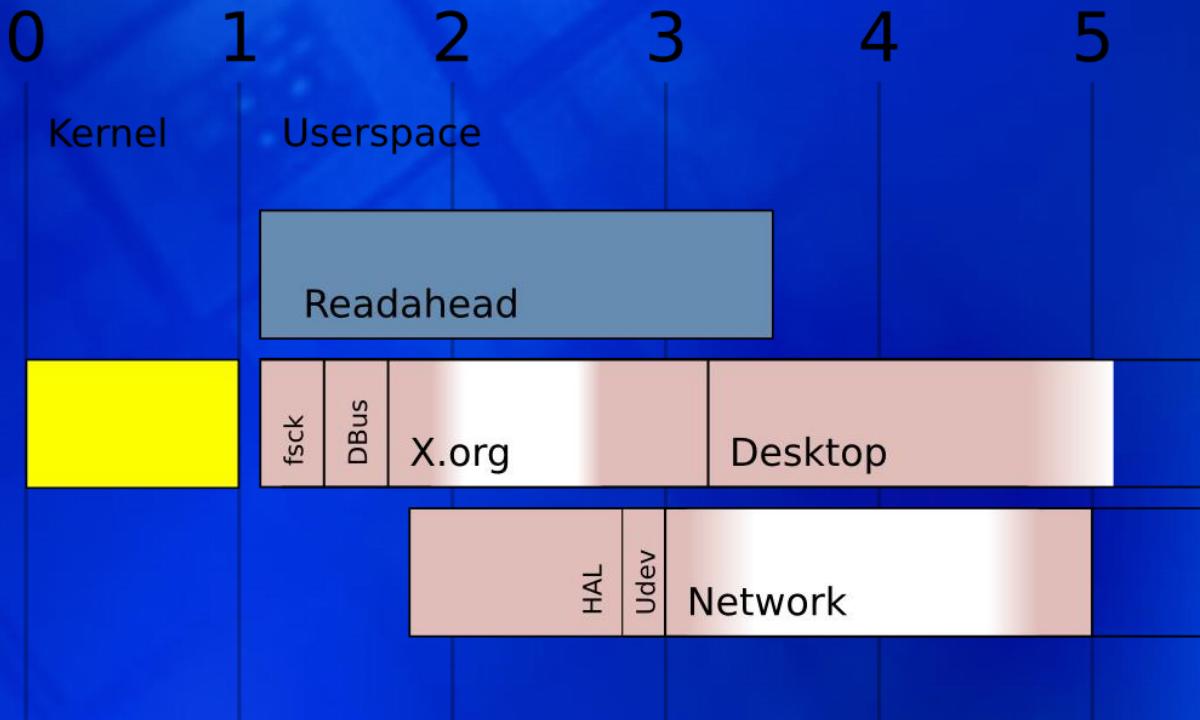
- ***sysvinit, not upstart: Asynchronous, not Parallel***
- **Udev**
 - **Persistent /dev reduces overhead enormously**



Early boot / Init

Budget: 1 second

- Asynchronous for non-critical path

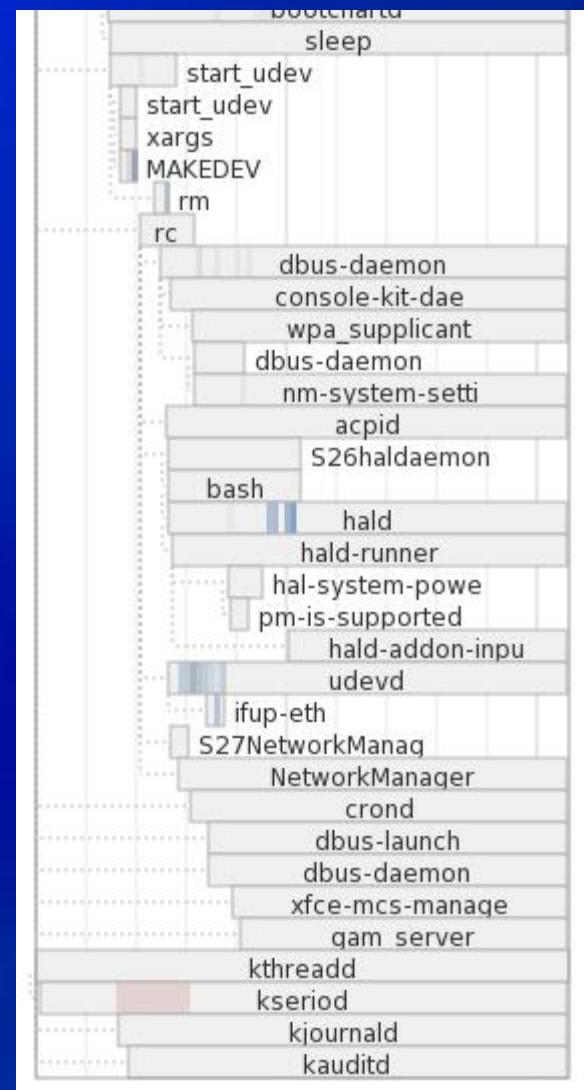
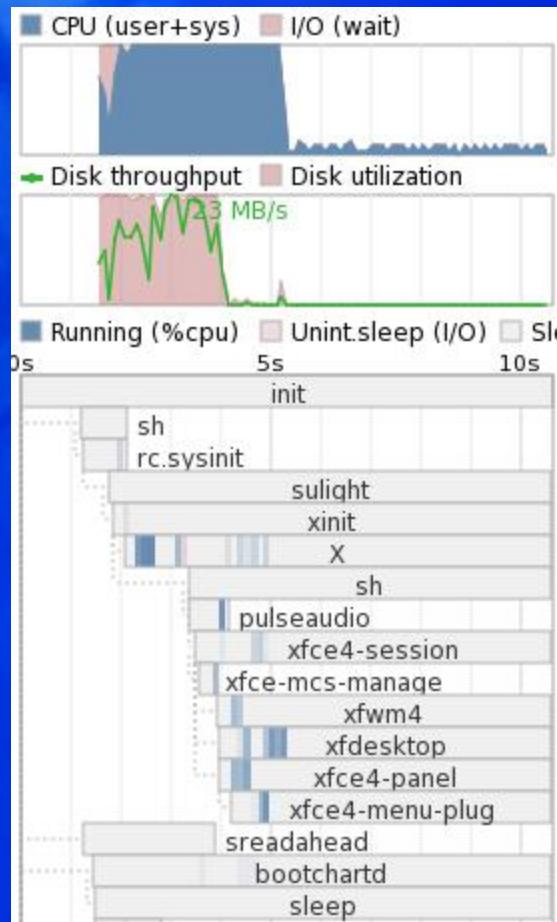
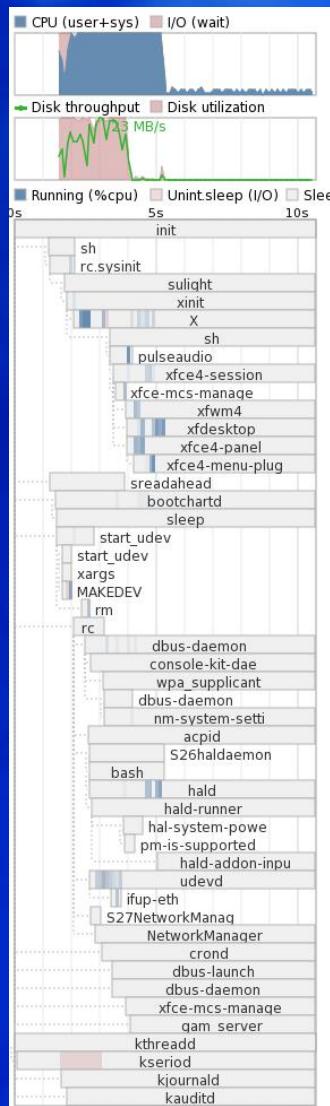


X.org

Budget: 1 second

- **xorg-x11-drv-intel driver**
 - Program everything during probing
 - ... then restore it to standard
 - ... then program it again for final use
 - Various "extra" delays trimmed (after fixing PCI posting bugs)
- **XKB**
 - Calling CPP ??? <explicit language removed>
 - Caching the result – Compute once, use forever





Fast suspend/resume

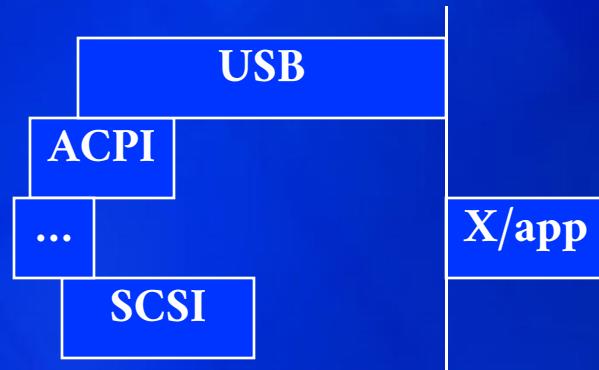
- Not as complex as fast boot
 - BIOS + kernel + X/application
 - No critical disk I/O request
- Time budget
 - Resume Kernel + X/application in 1 second

Driver optimization

- non-critical work can be deferred during device resume
 - Serio
 - offload resume to kseriod
 - Fix driver poor/wrong behavior
 - AHCI
 - HBA reset speed up

Device asynchronous resume

- Resume devices asynchronously
- Sync before resume X/applications



More aggressive attempts

- Device deferred resume instead of asynchronous resume
 - don't sync before resuming X/applications



- Screen is back before all the devices are resumed
- works but may bring potential risks

More aggressive attempts

- Resume X with the highest priorities
- Resume drm driver first to light the screen
- Resume X process before the other kernel drivers
- Works but not for upstream yet

S3 resume time Moblin

BIOS

kernel

X/app.

Current state

1.65s

1.78s

0.37s

Asynchronous resume

1.65s

1.19s

0.36s

deferred resume

1.65s

0.25s

0.36s

resume X first

1.65s

0.25s

0.20s

TODO

- push “deferred resume” mechanism upstream
- Fast suspend
- Fast shutdown

Acknowledgements

- Arjan van de Ven
- Auke Kok
- Li, Shaohua
- Zhao Yakui

