



CANONICAL

# Introduction to Runtime Power Management for IO devices

Ming Lei <ming.lei@canonical.com>

Kernel engineer, HWE team  
Canonical Ltd.

08 October 2011



# Outline

Introduction to Runtime PM

Runtime PM Framework

Runtime PM implementation examples

USB subsystem/OMAP3 or 4 platform

What is the next

References

# Introductions: Terms

- PM(power management): energy usage management
- System sleep: the whole system is put into sleep or hibernation
- RPM: Runtime Power Management
- Runtime PM: does not cover CPU, only for I/O devices in the presentation

# Introduction: Ideas behind Runtime PM

- Reduce the amount of energy used in whole system
- Limit the total power draw
- Not all devices are functional at most of time
- I/O devices which are in 'idle' state might be suspended
- Suspended I/O devices can be woken up by driver or itself (generally speaking, both are triggered by user) if they are to be used
- Similar in concepts to CPUIdle

# Introduction: Runtime PM overview

- Transparent to user
- Fine granularity PM, per device, not per whole system like system sleep
- Userspace application is not required to be frozen as system sleep
- For one device, runtime pm happens when the device is in 'idle' state and aren't going to be used in the near future, still very different with system sleep

# Introduction: Runtime PM overview

- What is the 'idle'?
  - Generally speaking, decided by device driver or its subsystem
  - For example:
    - Network interface: driver thinks it is idle after cable is unplugged
    - USB subsystem: no data transfer for some time
    - I2C: idle if no transfer is ongoing
    - ....
  - similar with CPUIdle, scheduler will check if cpu can be put in idle state
- BIOS independent(for x86)

# Introduction: Runtime PM overview

- Wakeup setting
  - remote wake-up may be enabled for runtime suspend but disallowed for system sleep
  - compared with system sleep, shorter wakeup latency is involved in runtime pm since system sleep always enters the lowest power consumption state but runtime pm doesn't

# Introduction: Why Do We Need a Framework for Device Runtime PM?

- System sleep is not enough to decrease runtime energy consumption
- Devices may depend on another device
- Be helpful to figure out 'idle' condition
- Not doable to do I/O runtime PM in CPUIdle
  - devices may be idle but cpu is not idle
  - task schedule may be caused during device suspend
  - some devices' suspend is moved from cpuidle platform driver to runtime PM (such as, uart suspend on omap3/4)
  - long latency is involved when returning from cpuidle
- 
-



# Introduction: Which subsystems/drivers has implemented runtime PM up to 3.1-rc4

- Usb subsystem  
HUB, usb mass storage, UVC, HID, CDC, serial, usb net, ...
- PCI subsystem  
e1000e, rtl8169, ehci-pci, uhci-pci, ...
- SCSI subsystem  
sd
- I2C subsystem
- MMC subsystem
- Serial devices
- Misc device(gpio, key,.....)
- ARCHs(RPM via dev\_pm\_domain)  
arch/sh/kernel/cpu/shmobile/pm\_runtime.c  
arch/arm/mach-omap1/pm\_bus.c  
arch/arm/mach-shmobile/pm\_runtime.c  
arch/arm/plat-omap/omap\_device.c
-



# Runtime PM Framework

# Runtime PM framework

- Device “States”

Runtime PM framework uses abstract states of devices

- ACTIVE: Device can do I/O (presumably in the full-power state).
- SUSPENDED: Device cannot do I/O (presumably in a low-power state).
- SUSPENDING: Device state is changing from ACTIVE to SUSPENDED.
- RESUMING: Device state is changing from SUSPENDED to ACTIVE.

# Runtime PM framework

- callbacks to be implemented by subsystem or driver

include/linux/pm.h

```
struct dev_pm_ops {  
    ...  
    int (*runtime_suspend)(struct device *dev);  
    int (*runtime_resume)(struct device *dev);  
    int (*runtime_idle)(struct device *dev);  
};
```

- RESUMINGING → ACTIVE: ->runtime\_resume succeed
- RESUMINGING → SUSPENDED: ->runtime\_resume fails
- SUSPENDING → SUSPENDED: ->runtime\_suspend succeed
- SUSPENDING → ACTIVE: ->runtime\_suspend fails

# Runtime PM framework

- ->runtime\_idle
  - check if the device is idle and can be suspended
  - if yes, call appropriate runtime PM API to [schedule/queue] auto suspend/suspend for the device
- ->runtime\_suspend
  - configure remote wakeup signal if it is allowed and capable
  - execute platform specific things(clock, power,...) to suspend the device
- ->runtime\_resume
  - execute platform specific things(clock, power,...) to wakeup the device

# Runtime PM framework

- Where to implement callbacks and its priority

XX: suspend/resume/idle

if (dev->pm\_domain)

    callback = dev->pm\_domain->ops.runtime\_XX;

else if (dev->type && dev->type->pm)

    callback = dev->type->pm->runtime\_XX;

else if (dev->class && dev->class->pm)

    callback = dev->class->pm->runtime\_XX;

else if (dev->bus && dev->bus->pm)

    callback = dev->bus->pm->runtime\_XX;

else

    callback = NULL;

# Runtime PM framework: Callback

- `irq_safe` is set if the callbacks can be run in interrupt off contexts
- `irq_safe` is not set  
    `dev->power.lock` is released and irq is enabled before callbacks are called
- `irq_safe` is set  
    `dev->power.lock` is released but irq is not enabled before callbacks are called
- Release of power lock is the root of all kinds of races

# Runtime PM framework: API

- Suspend related API
  - `pm_runtime_suspend`
  - `pm_schedule_suspend`
  - `pm_runtime_autosuspend`
  - `pm_request_autosuspend`
- Resume related API
  - `pm_runtime_resume`
  - `pm_request_resume`
- Idle related API
  - `pm_runtime_idle`
  - `pm_request_idle`



# Runtime PM framework: API

- Reference Counting

Device with references held cannot be suspended/Idled

Device can be resumed no matter if its reference is zero or not

- Get reference

`pm_runtime_get`

`pm_runtime_get_sync`

`pm_runtime_get_noresume`

- Put reference

`pm_runtime_put`

`pm_runtime_put_sync`

`pm_runtime_put_autosuspend`

`pm_runtime_put_sync_suspend`

`pm_runtime_put_sync_autosuspend`

`pm_runtime_put_noidle`

# Runtime PM framework: API

- Misc API

`pm_runtime_enable/pm_runtime_disable`

`pm_runtime_allow/pm_runtime_forbid`

`pm_runtime_irq_safe`

`pm_suspend_ignore_children`

`pm_runtime_barrier`

`device_run_wake/device_set_run_wake`

`pm_runtime_use_autosuspend`

`pm_runtime_set_autosuspend_delay`

`pm_runtime_autosuspend_expiration`

`pm_runtime_mark_last_busy`

# RPM framework: Idle notification

- Put reference and check if the usage count is zero

`pm_runtime_put`

`pm_runtime_put_sync`

- After `->runtime_resume` successfully  
queue a idle notification to the device
- After `->runtime_suspend` successfully  
queue a idle notification to its parent
- Before allowing autosuspend

# RPM framework: Auto Suspend

- First introduced in usb subsystem
- User can set 'autosuspend delay' or use the default value
- `pm_runtime_mark_last_busy` marks the last time of the device's busy state
- Once auto suspend is scheduled or executed, runtime PM will
  - check if the difference between current time and last busy is beyond 'autosuspend delay' of the device.
  - If yes, `->runtime_suspend` will be called for the device

# RPM framework: Interaction between RPM and system sleep

- Device is suspended before system sleep
  - may runtime resume first in .suspend for different remote wakeup settings
  - During system resume, the simplest approach is to bring all devices back to full power, even if they had been suspended before the system suspend began
- During system suspend
  - pm\_runtime\_get\_noresume(dev)
  - pm\_runtime\_barrier(dev)
  - abort suspend if wakeup events are found
  - call ->suspend(dev)
  - pm\_runtime\_disable(dev)
- During system resume
  - pm\_runtime\_enable(dev)
  - call ->resume(dev)
  - pm\_runtime\_put\_sync(dev)

# RPM framework: Wakeup signal

- Depend on the platform and bus type
  - Special signals from low-power states (device signal causes another device to generate an interrupt).  
PCI Power Management Event (PME) signals.  
USB "remote wakeup".
  - Interrupts from low-power states (wakeup interrupts).
- What is needed?
  - Power Domain/Subsystem and/or driver callbacks need to set up devices to generate these signals.
  - The resulting interrupts need to be handled (devices should be put into the ACTIVE state as a result).



# Runtime PM implementation examples USB subsystem / OMAP3 or 4

# RPM implementation in usb subsystem

- Usb subsystem implementation(drivers/usb/core/driver.c)

->runtime\_suspend

```
if (autosuspend_check(udev) != 0)
```

```
    return -EAGAIN;
```

```
status = usb_suspend_both(udev, PMSG_AUTO_SUSPEND);
```

->runtime\_resume

```
status = usb_resume_both(udev, PMSG_AUTO_RESUME);
```

->runtime\_idle

```
if (autosuspend_check(udev) == 0)
```

```
    pm_runtime_autosuspend(dev);
```



# RPM implementation in usb subsystem

- Usb interface driver
  - simple usage(UVC class, usb-skeleton, ...)
    - resume during ->open
    - suspend during ->release
  - complicated examples(cdc-acm, cdc\_ether, HID, ...)
    - wakeup device only during urb transfer

# RPM implementation in OMAP3/4 platform

- Platform bus
- Power domain(`arch/arm/plat-omap/omap_device.c`)

```
static struct dev_pm_domain omap_device_pm_domain = {  
    .ops = {  
        SET_RUNTIME_PM_OPS(_od_runtime_suspend, _od_runtime_resume,  
                            _od_runtime_idle)  
        USE_PLATFORM_PM_SLEEP_OPS  
        .suspend_noirq = _od_suspend_noirq,  
        .resume_noirq = _od_resume_noirq,  
    }  
};
```

# RPM implementation in OMAP3/4 platform

- omap\_device\_register

```
od->pdev.dev.parent = &omap_device_parent;
od->pdev.dev.pm_domain = &omap_device_pm_domain;
return platform_device_register(&od->pdev);
```

- \_od\_runtime\_idle(struct device \*dev)

```
->pm_generic_runtime_idle(dev);
    ->dev->driver->pm->runtime_idle(dev)    /*if defined*/
```

- \_od\_runtime\_suspend(struct device \*dev)

```
->pm_generic_runtime_suspend(dev);
    ->dev->driver->pm->runtime_suspend(dev)    /*if defined*/
->omap_device_idle(pdev); /*configure idle mode, disable clocks,...*/
```

# RPM implementation in OMAP3/4 platform

- `_od_runtime_resume(struct device *dev)`

```
->omap_device_enable(pdev);    /*configure idle mode, enable clocks,...*/  
->pm_generic_runtime_resume(dev);  
->->dev->driver->pm->runtime_resume(dev)    /*if defined*/
```

- `omap_device(arch/arm/plat-omap/include/plat/omap_device.h)`

```
struct omap_device {  
    struct platform_device    pdev;  
    struct omap_hwmod    **hwmods;                /*IPs inside the device*/  
    u8    hwmods_cnt;  
    struct omap_device_pm_latency    *pm_lats;        /*multiple latency level*/  
    u8    pm_lats_cnt;  
    .....  
}
```

# RPM implementation in OMAP3/4 platform

- arch/arm/plat-omap/include/plat/omap\_hwmod.h

```
-struct omap_hwmod {  
    .....  
    struct clk      _clk;                /*function clock*/  
    struct clockdomain *clkdm;           /*clock domain of the function clock*/  
    struct omap_hwmod_ocp_if **masters;  /*interfaces that this hwmod can initiate on*/  
    struct omap_hwmod_ocp_if **slaves;   /*interfaces that this hwmod can respond on*/  
    .....  
}  
  
struct omap_hwmod_ocp_if {  
    struct omap_hwmod *master;           /*initiator*/  
    struct omap_hwmod *slave;            /*target*/  
    struct clk      *_clk;               /*interface clock*/  
    .....  
}
```

# RPM implementation in OMAP3/4 platform

- more fine granularity PM unit
  - hwmod is the basic PM unit instead of device
- different view about PM dependency compared with device model(parent vs. child)
  - more hwmods are involved for one device's PM
- Universal clock/clockdomain/.. hardware handling during suspend/resume
- Support multiple latency level



# What is the next

# What is the next

- Implement RPM in more subsystems or drivers
- PM Qos per device
  - patches from TI will enter 3.2
  - wakeup latency request can be put on per device
  - user space API is under discussion
- PM dependency out of device model
- No idle knowledge in some devices' drivers
  - user space hint is needed?
  - Or device specific implementation





# References

# References

- `drivers/base/power/`
- `drivers/usb/core/`
- `arch/arm/plat-omap/omap_device.c`
- `arch/arm/mach-omap2/`
- `Include/linux/pm*.h`
- `Documents/power/runtime_pm.txt`
- R. J. Wysocki, Runtime Power Management Framework for I/O Devices in the Linux Kernel  
([http://events.linuxfoundation.org/slides/2010/linuxcon2010\\_wysocki.pdf](http://events.linuxfoundation.org/slides/2010/linuxcon2010_wysocki.pdf)).



CANONICAL

Questions please  
Thank you

