



PerfEvent 与 Intel PMU 介绍

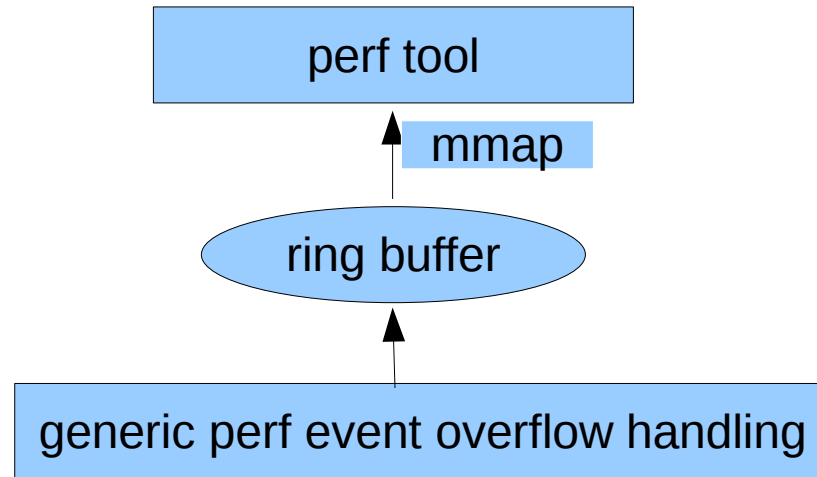
林铭 Intel 开源技术中心
ming.m.lin@intel.com

PerfEvent 简介



- 功能强大的性能分析工具
 - kernel code + userspace tool
 - 支持多种 CPU: x86, PowerPC, ARM, Sparc, MIPS, Alpha, Blackfin, SH ...
- 范围
 - 整个系统 (所有 CPU + 所有进程)
 - 某个或某几个 CPU
 - 某个进程
- 对象
 - hardware event: cpu cycles, instructions, cache hit/miss
 - software event: page fault, context switch, cpu migrations
 - tracepoint
- 形式
 - 采样 (实时观测, 保存分析)
 - 统计
- 收集 call graph
- Kernel/userspace 分别收集

PerfEvent 架构



用户运行 perf tool:
perf top
perf record
perf stat

hardware events

NMI interrupt handler

Hardware counter
overflow interrupt

software events

perf_sw_event()

perf_swevent_hrtimer

tracepoint

Intel PMU



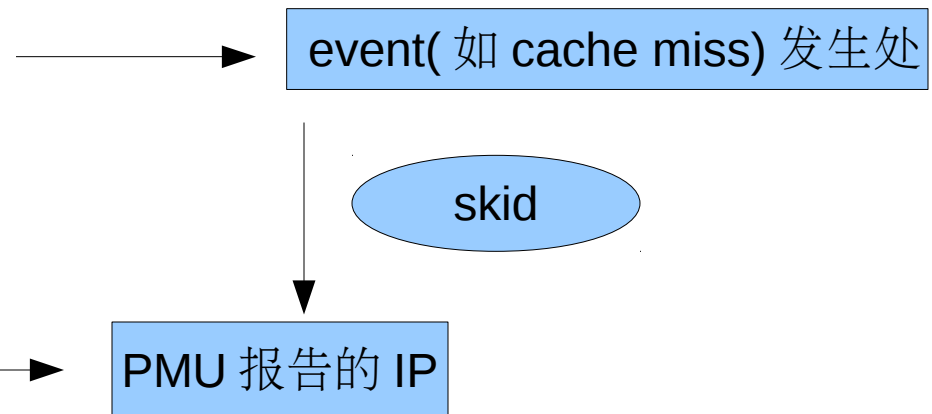
- PMU: Performance Monitoring Unit
- performance counter 和 performance event
 - counter 记录 event 发生的次数
 - counter: Core2 2 个, Nehalem/Westmere/SandyBridge: 4 个
 - event: 架构化 7 个: cpu cycles, bus cycles, cpu instruction, LLC hit, LLC miss, branch instruction, branch misses, 非架构化上百个
 - counter overflow interrupt
- PerfEvent counter 调度

skid issue



Skid: PMU 报告的 IP 不是 event 发生的准确 IP

```
400400: push %rbp
400401: mov  %rsp,%rbp
400404: push %rbx
400405: sub  $0x8,%rsp
400409: cmpb $0x0,0x200420(%rip)
400410: jne  40045d
400412: mov  $0x600648,%ebx
400417: mov  0x20041a(%rip),%rax
40041e: sub  $0x600640,%rbx
400425: sar  $0x3,%rbx
400429: sub  $0x1,%rbx
40042d: cmp  %rbx,%rax
400430: jae  400456
400432: nopw 0x0(%rax,%rax,1)
```



PEBS(Precise Event Based Sampling)



- PEBS record format

Table 30-12. PEBS Record Format for Intel Core i7 Processor Family

Byte Offset	Field	Byte Offset	Field
0x0	R/EFLAGS	0x58	R9
Byte Offset	Field	Byte Offset	Field
0x8	R/EIP	0x60	R10
0x10	R/EAX	0x68	R11
0x18	R/EBX	0x70	R12
0x20	R/ECX	0x78	R13
0x28	R/EDX	0x80	R14
0x30	R/ESI	0x88	R15
0x38	R/EDI	0x90	IA32_PERF_GLOBAL_STATUS
0x40	R/EBP	0x98	Data Linear Address
0x48	R/ESP	0xA0	Data Source Encoding
0x50	R8	0xA8	Latency value (core cycles)

- PEBS events

- cycles, instructions, branch instructions, branch misses

- PEBS 用法示例: `perf top -e cycles:p`

PEBS IP+1 issue



```
400400: push  %rbp
400401: mov   %rsp,%rbp
400404: push  %rbx
400405: sub   $0x8,%rsp
400409: cmpb  $0x0,0x200420(%rip)
400410: jne   40045d
400412: mov   $0x600648,%ebx
400417: mov   0x20041a(%rip),%rax
40041e: sub   $0x600640,%rbx
400425: sar   $0x3,%rbx
400429: sub   $0x1,%rbx
40042d: cmp   %rbx,%rax
400430: jae   400456
400432: nopw  0x0(%rax,%rax,1)
```

event 发生处

PEBS 报告的 IP

LBR (Last Branch Record)



- LBR MSR

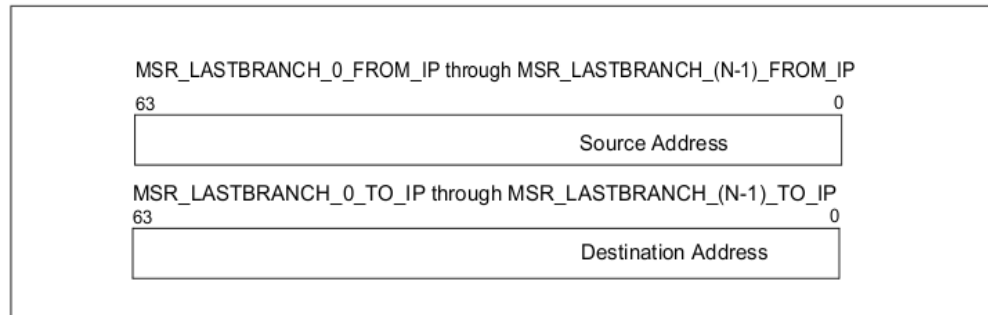


Figure 16-4. 64-bit Address Layout of LBR MSR

```
4004ec:    movl    $0x0,-0x4(%rbp)
4004f3:    cmpl    $0x0,-0x4(%rbp)
4004f7:    je      400505
4004f9:    mov     $0x40060c,%edi
4004fe:    callq   4003e0
400503:    jmp     40050f
400505:    mov     $0x40060c,%edi
40050a:    callq   4003e0
40050f:    mov     $0x0,%eax
```

Source Address: 4004f7
Destination Address: 400505

LBR: PEBS IP+1 issue fixup



- 第一种情况: event 发生在 branch instruction 处

```
4004f7:    je     400505    —————→ event 发生处
4004f9:    mov    $0x40060c,%edi
4004fe:    callq  4003e0
400503:    jmp     40050f
400505:    mov    $0x40060c,%edi    —→ PEBS 记录的 IP
40050a:    callq  4003e0
```

Source Address: 4004f7
Destination Address: 400505

将 PEBS IP 纠正为 Source Address
用法示例: `perf top -e cycles:pp`

LBR: PEBS IP+1 issue fixup cont.



- 第二种情况: event 不是发生在 branch instruction 处

```
400503:    jmp     40050f
400505:    mov     $0x40062c,%edi
40050a:    callq   4003e0
40050f:    addl     $0x1,-0x4(%rbp)
400513:    movl     $0x64,-0x4(%rbp)
40051a:    mov     -0x4(%rbp),%ecx
40051d:    mov     $0x66666667,%edx
```

event 发生处

PEBS 记录的 IP

Source Address: 400503
Destination Address: 40050f

将 PEBS IP 纠正为 40051a

Kernel instruction decoder:

40050f: 83 45 fc 01 (长度 4)

400513: c7 45 fc 64 00 00 00 (长度 8)

40051a: 8b 4d fc (长度 3)

40051d: ba 67 66 66 66 (长度 5)

LBR: sampling taken branches



```
void f2(void) {}
void f3(void) {}
void f1(unsigned long n)
{
    if (n & 1UL)
        f2();
    else
        f3();
}
int main(void)
{
    unsigned long i;
    for (i=0; i < N; i++)
        f1(i);
    return 0;
}
```

```
$ perf record -b any_call -e cycles:u branchy
$ perf report
# Events: 19K cycles
#
# Overhead Source Symbol Target Symbol
# .....
#
52.50%    [.] main          [.] f1
23.99%    [.] f1          [.] f3
23.48%    [.] f1          [.] f2
```

Patches by Stephane Eranian @google
[PATCH 00/12] perf_events: add support for
sampling taken branches
<http://marc.info/?l=linux-kernel&m=131805702307992&w=2>

Load latency & Precise store



- Load latency
 - Load data linear address
 - Latency value
 - Data source
- Precise store
 - Store data linear address
 - Store status
- Patches

[PATCH 0/4] perf: memory load/store events generalization

<http://marc.info/?l=linux-kernel&m=130976621805927&w=2>

perf mem 例子 1



```
$ perf mem
```

```
usage: perf mem [<options>] {record <command> |report}
```

```
-t, --type <type>      memory operations(load/store)
-L, --latency <n>       latency to sample(only for load op)
```

```
$ perf mem -t load record make -j8
```

```
<building kernel ..., monitoring memory load operation>
```

```
$ perf mem -t load report
```

```
Memory load operation statistics
```

```
=====
L1-local: total latency= 28027, count= 3355(avg=8)
L2-snoop: total latency= 1430, count= 29(avg=49)
L2-local: total latency= 124, count= 8(avg=15)
L3-snoop, found M: total latency= 452, count= 4(avg=113)
L3-snoop, found no M: total latency= 0, count= 0(avg=0)
L3-snoop, no coherency actions: total latency= 875, count= 18(avg=48)
L3-miss, snoop, shared: total latency= 0, count= 0(avg=0)
L3-miss, local, exclusive: total latency= 0, count= 0(avg=0)
L3-miss, local, shared: total latency= 0, count= 0(avg=0)
L3-miss, remote, exclusive: total latency= 0, count= 0(avg=0)
L3-miss, remote, shared: total latency= 0, count= 0(avg=0)
Unknown L3: total latency= 0, count= 0(avg=0)
IO: total latency= 0, count= 0(avg=0)
Uncached: total latency= 464, count= 30(avg=15)
```

```
$ perf mem -t store record make -j8
```

```
<building kernel ..., monitoring memory store operation>
```

```
$ perf mem -t store report
```

```
Memory store operation statistics
```

```
=====
data-cache hit: 8138
data-cache miss: 0
STLB hit: 8138
STLB miss: 0
Locked access: 0
Unlocked access: 8138
```

perf mem 例子 2



Data structure 分析

```
    struct bar {  
0.00 :      int poekoe[5];  
60.00 :      int fubar;  
    };  
  
    struct foo {  
0.00 :      long poekoe[3];  
10.00 :      struct bar *bar;  
    };  
  
    struct tmp {  
0.00 :      long poekoe[4];  
30.00 :      int blah;  
    };
```

将地址与 structure field 关联

Questions?



Legal Information



- INFORMATION IN THIS DOCUMENT IS PROVIDED IN CONNECTION WITH INTEL[™] PRODUCTS. EXCEPT AS PROVIDED IN INTEL'S TERMS AND CONDITIONS OF SALE FOR SUCH PRODUCTS, INTEL ASSUMES NO LIABILITY WHATSOEVER, AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY RELATING TO SALE AND/OR USE OF INTEL PRODUCTS, INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT, OR OTHER INTELLECTUAL PROPERTY RIGHT.
- Intel may make changes to specifications, product descriptions, and plans at any time, without notice.
- All dates provided are subject to change without notice.
- Intel and Intel[™] Core[™] i7 are trademarks of Intel Corporation in the U.S. and other countries.
- Other names and brands may be claimed as the property of others.
- Copyright © 2011 Intel Corporation. All rights reserved.