

# On The Way to a Healthy Btrfs Towards Enterprise

Miao Xie

<miaox@cn.fujitsu.com>

# Tagged with “experimental”

- Oops still remains some.
- Fsck is at experimental stage.
- Immature for production use.

# Our goal

- Improve btrfs to be fit for production use.

# Btrfs that enterprise requests for

- Good performance
- Reliability
- Scalability
- Fault tolerance
- Features
  - Snapshot
  - Integrated multiple devices support
  - transparent compression
  - etc


# Our progress

- Error handling infrastructure
- Free space cluster per node
- Snapshot aware defrag
- Inode cache
- Per file cow and compression control
- Extent buffer cache
- Rbtree lock contention
- A great amount of bug fixes and cleanups

# Progress:

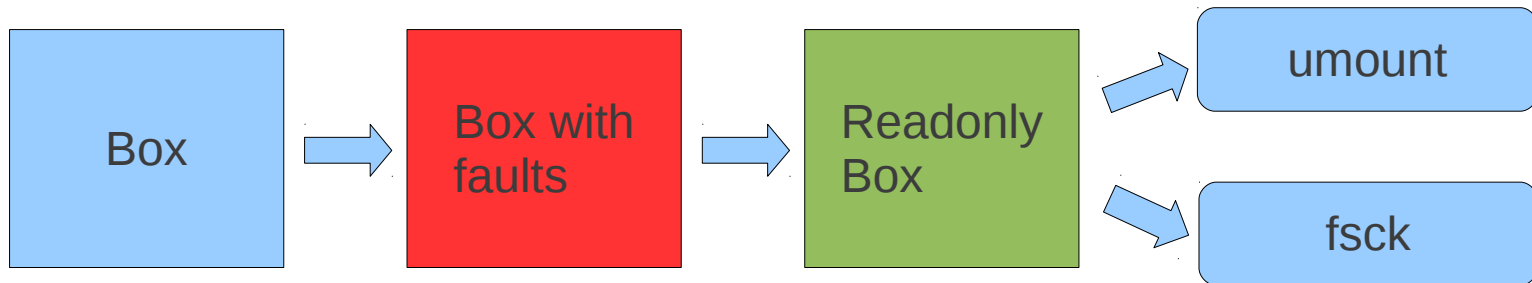
## Error handling infrastructure

- Forced readonly mounts on errors
  - Being fault tolerant of disk corruptions
  - Build a framework which can flip btrfs to readonly when there are errors
  - Replace BUG() and BUG\_ON()



w/o patch

w patch



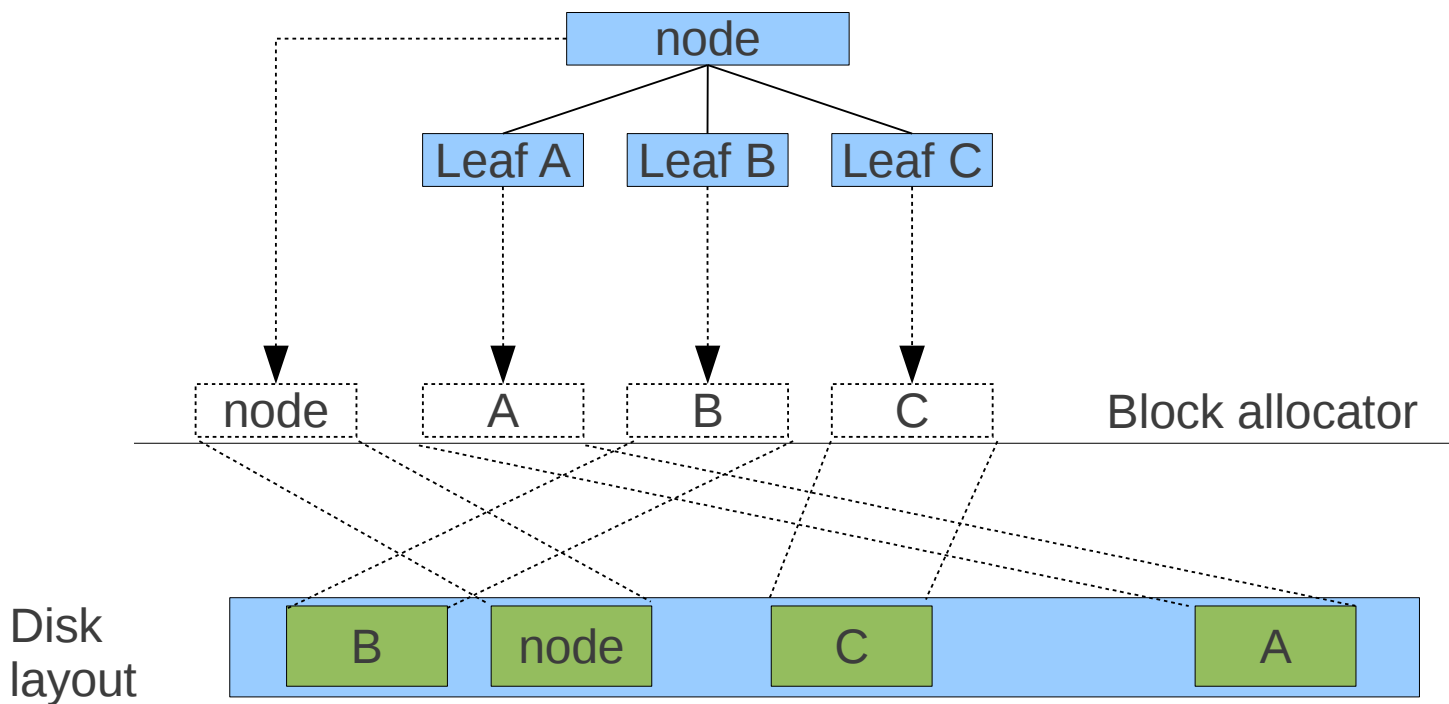
# Progress:

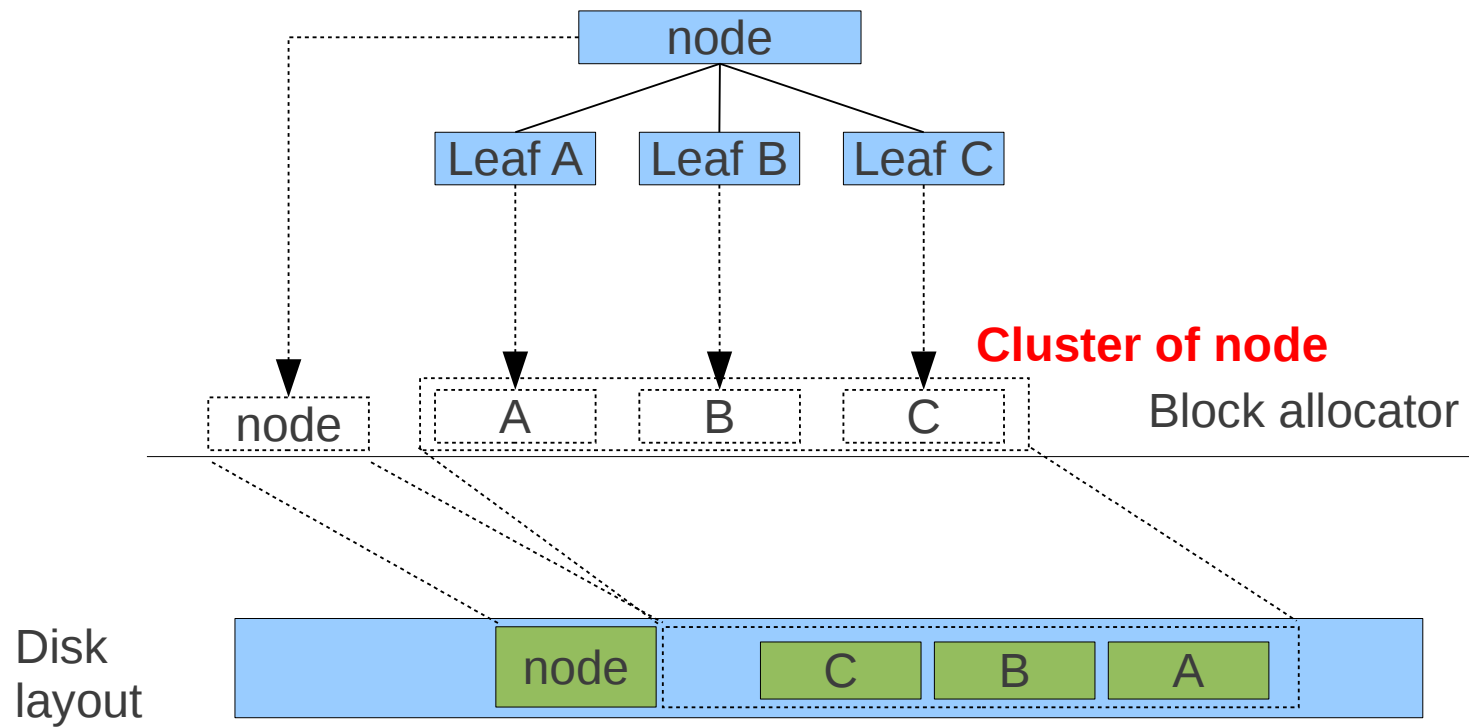
## free space cluster per node



- Reduce metadata fragments
- Improve (small files) sequential read performance



# WAFL(Write Anywhere File Layout)

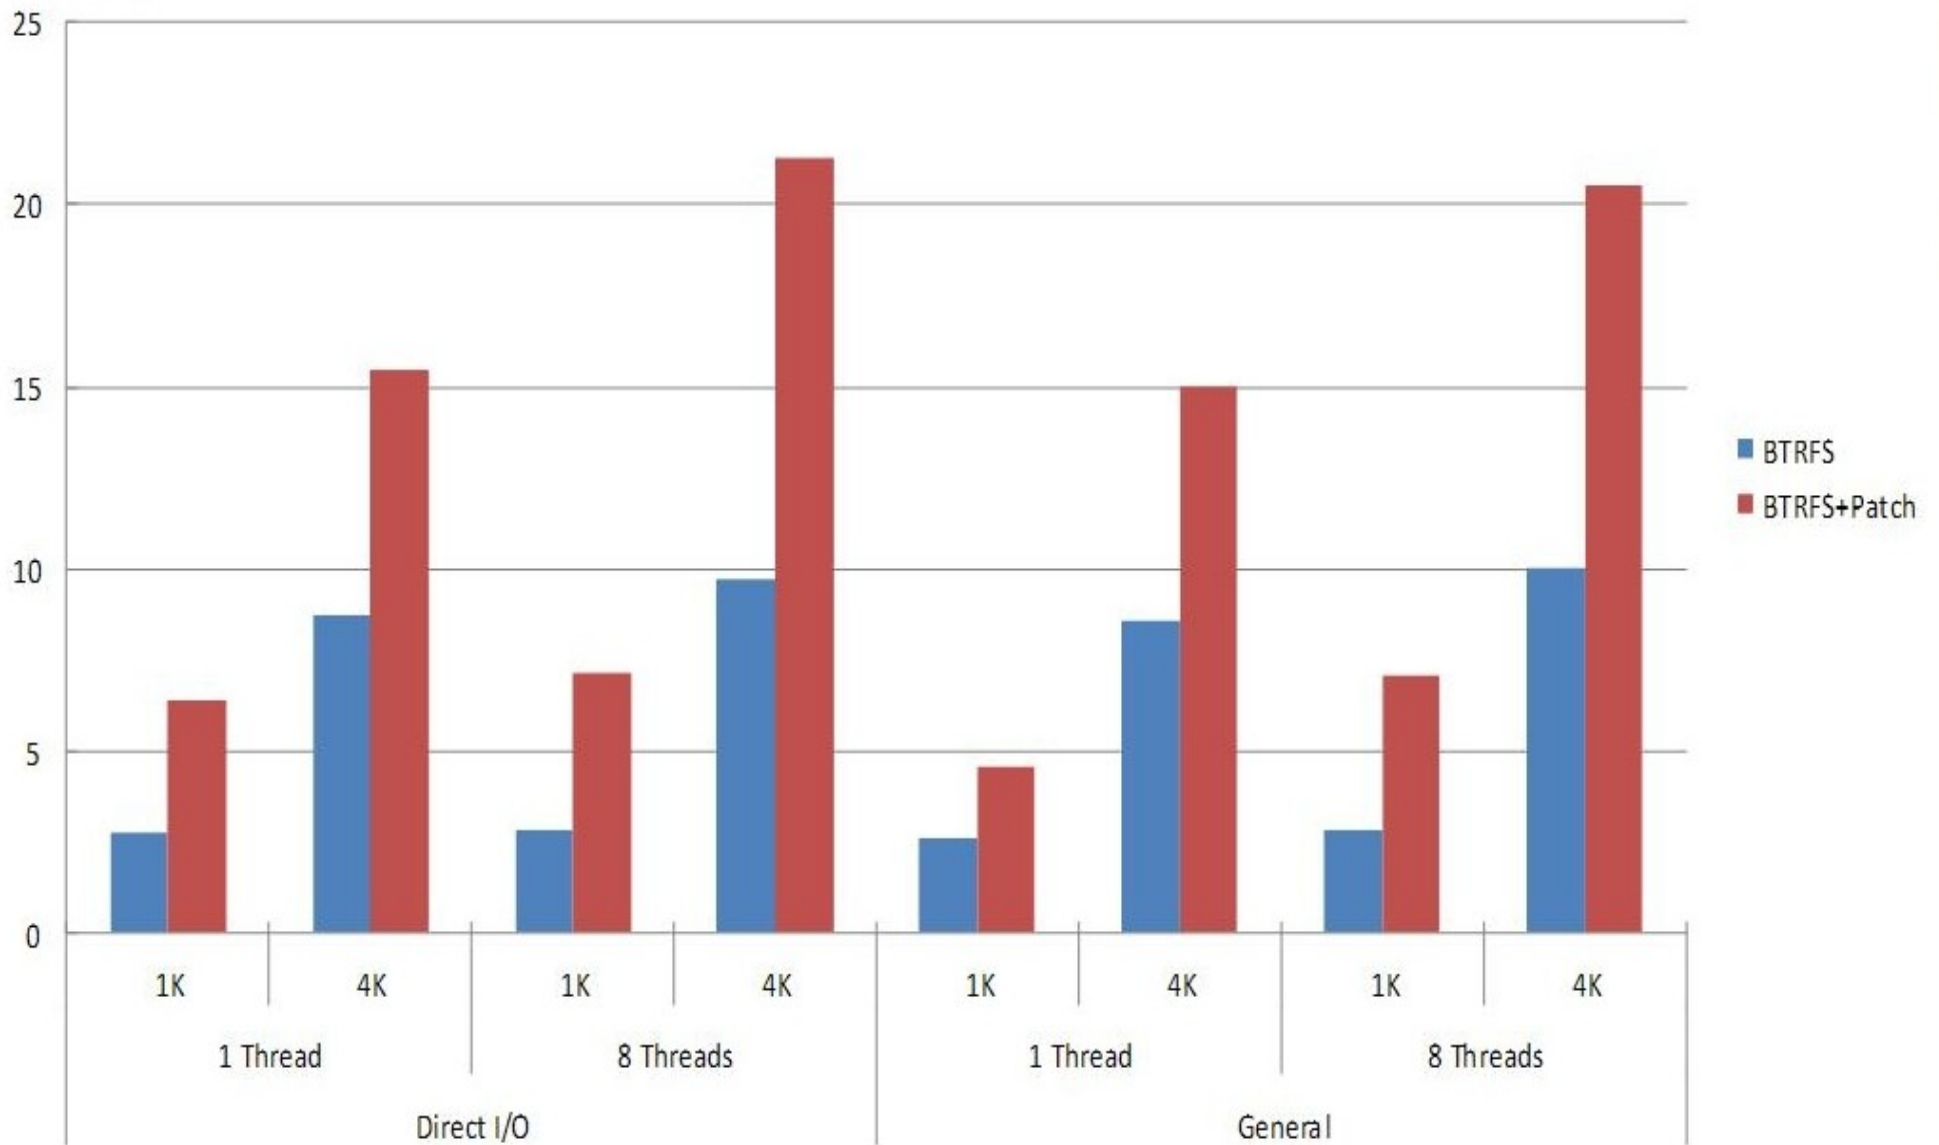






## Small files sequential read (Based on 2.6.38)

Unit: MB/s

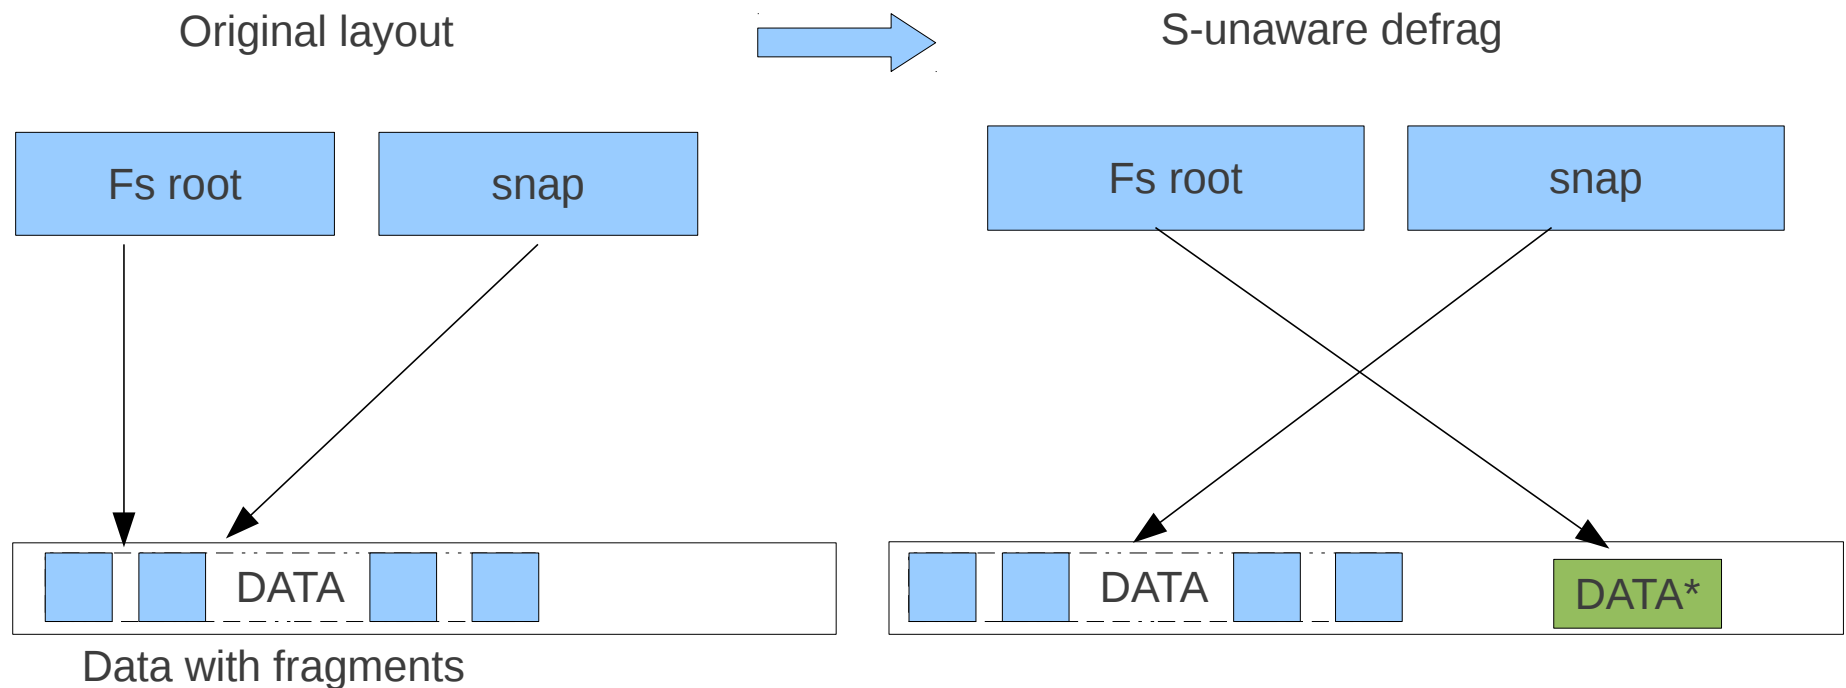


# Progress:

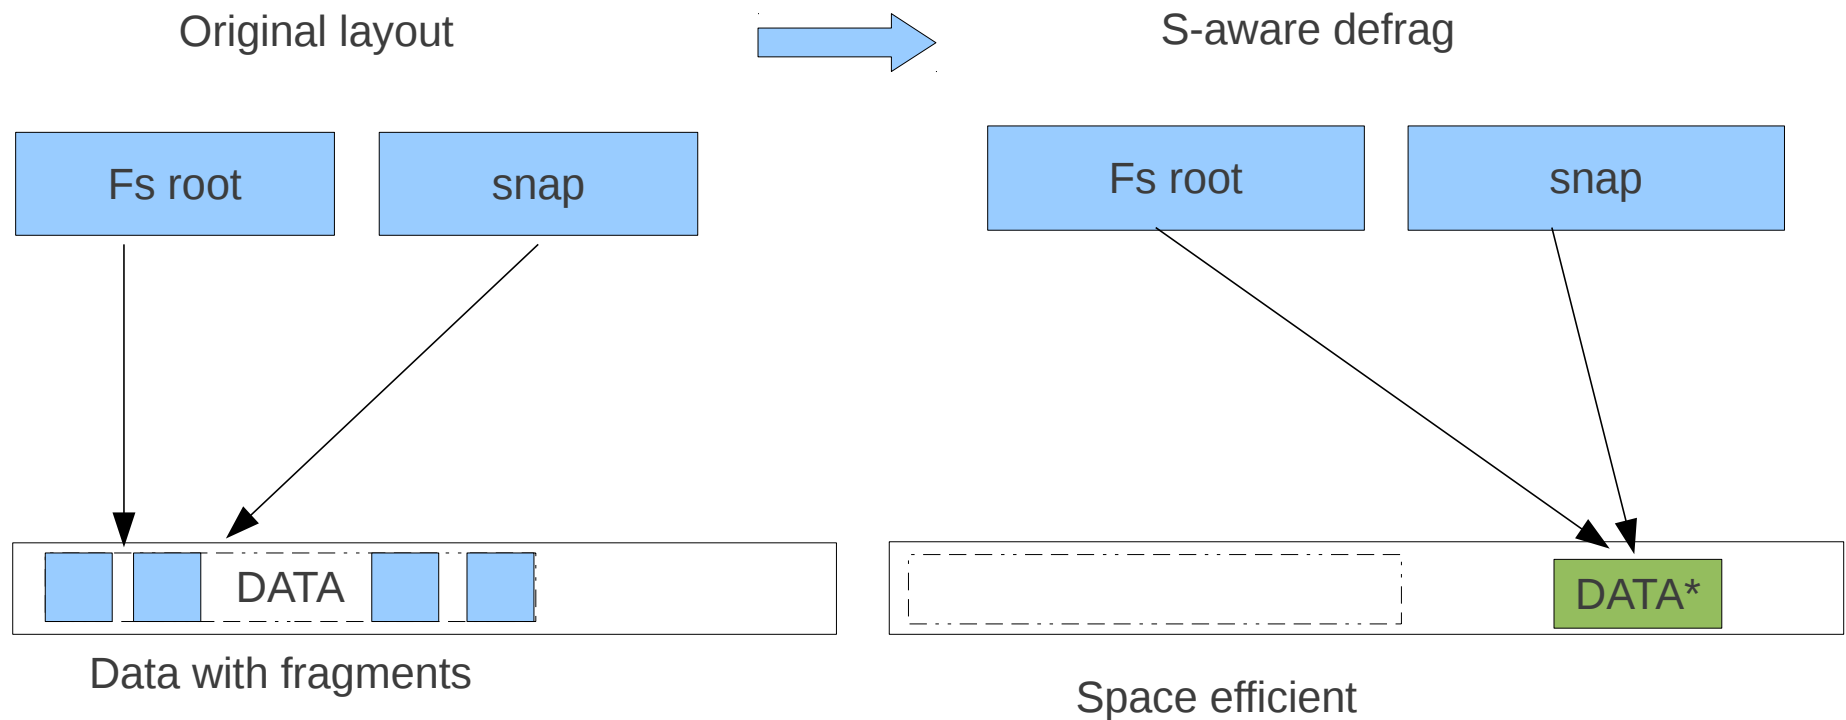
## snapshot aware defrag

- Make defragment code preserve the sharing among snapshots.
  - Btrfs has designed **back references** for this

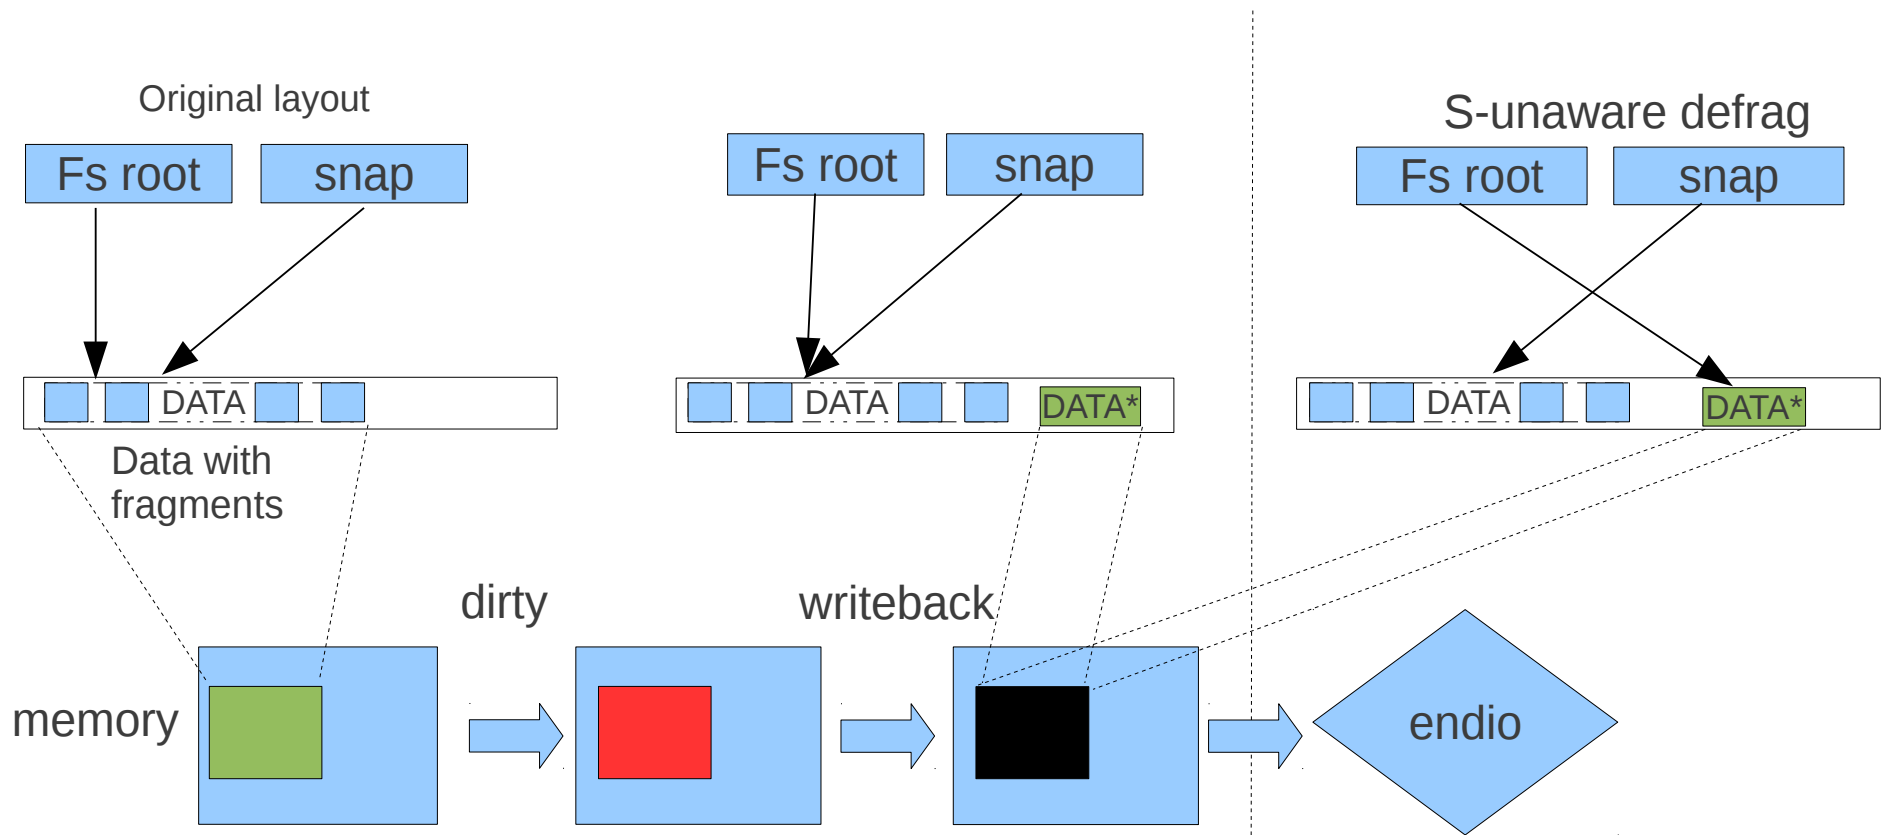
# Snapshot unaware defragment



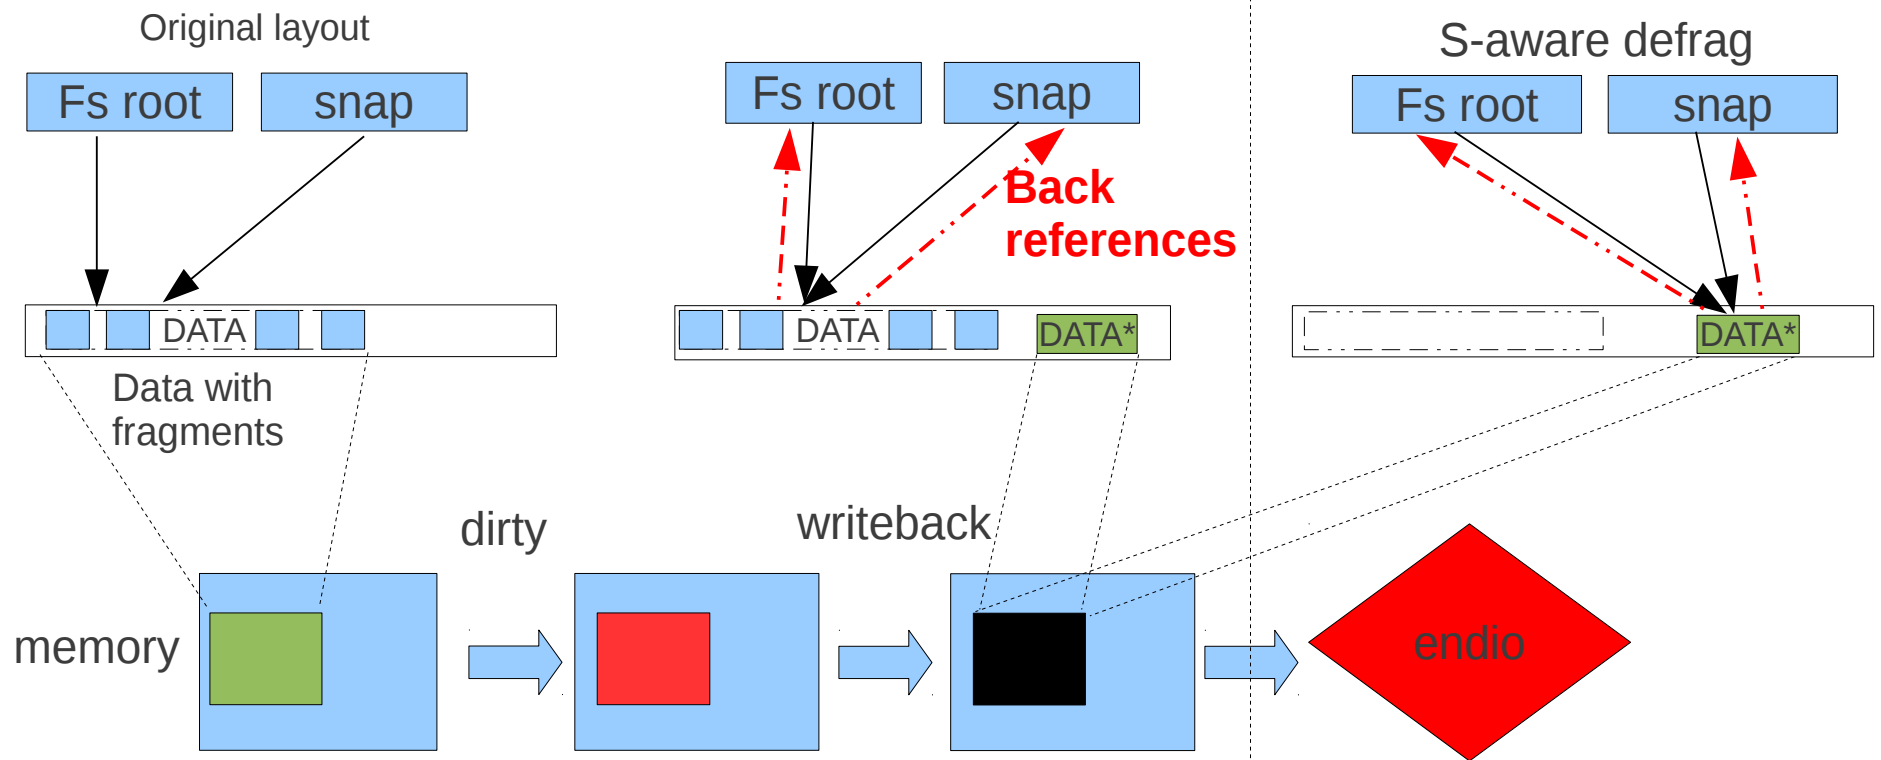
# Snapshot aware defragment



# How does defragment work









# Progress: inode cache


- Without inode cache,
  - Will not reclaim inode number when deleting files
  - It will not reuse inode number




# Progress:

## Per file cow and compression control

- Beside mount options, we need to control these flags on a per-inode basis.




Mount -o compress  
Mount -o nodatacow




w/o patch

w patch



Chattr -C E  
Chattr -c E




# Progress:

## extent buffer cache

- Extent buffer is a basic unit of metadata
- Expensive on searching and reading
- Cache misses depend on workloads

# What is extent buffer





# Progress:

## rbtree lock contention

- Why
  - Lock contention is really critical on performance
- How
  - Some rbtrees are domained by reads
  - Lockless read
- What
  - Build 'read mostly' circumstance
    - Find where the write locks are held
    - Try to reduce them as much as possible
  - Apply RCU, or rwlock






extent state

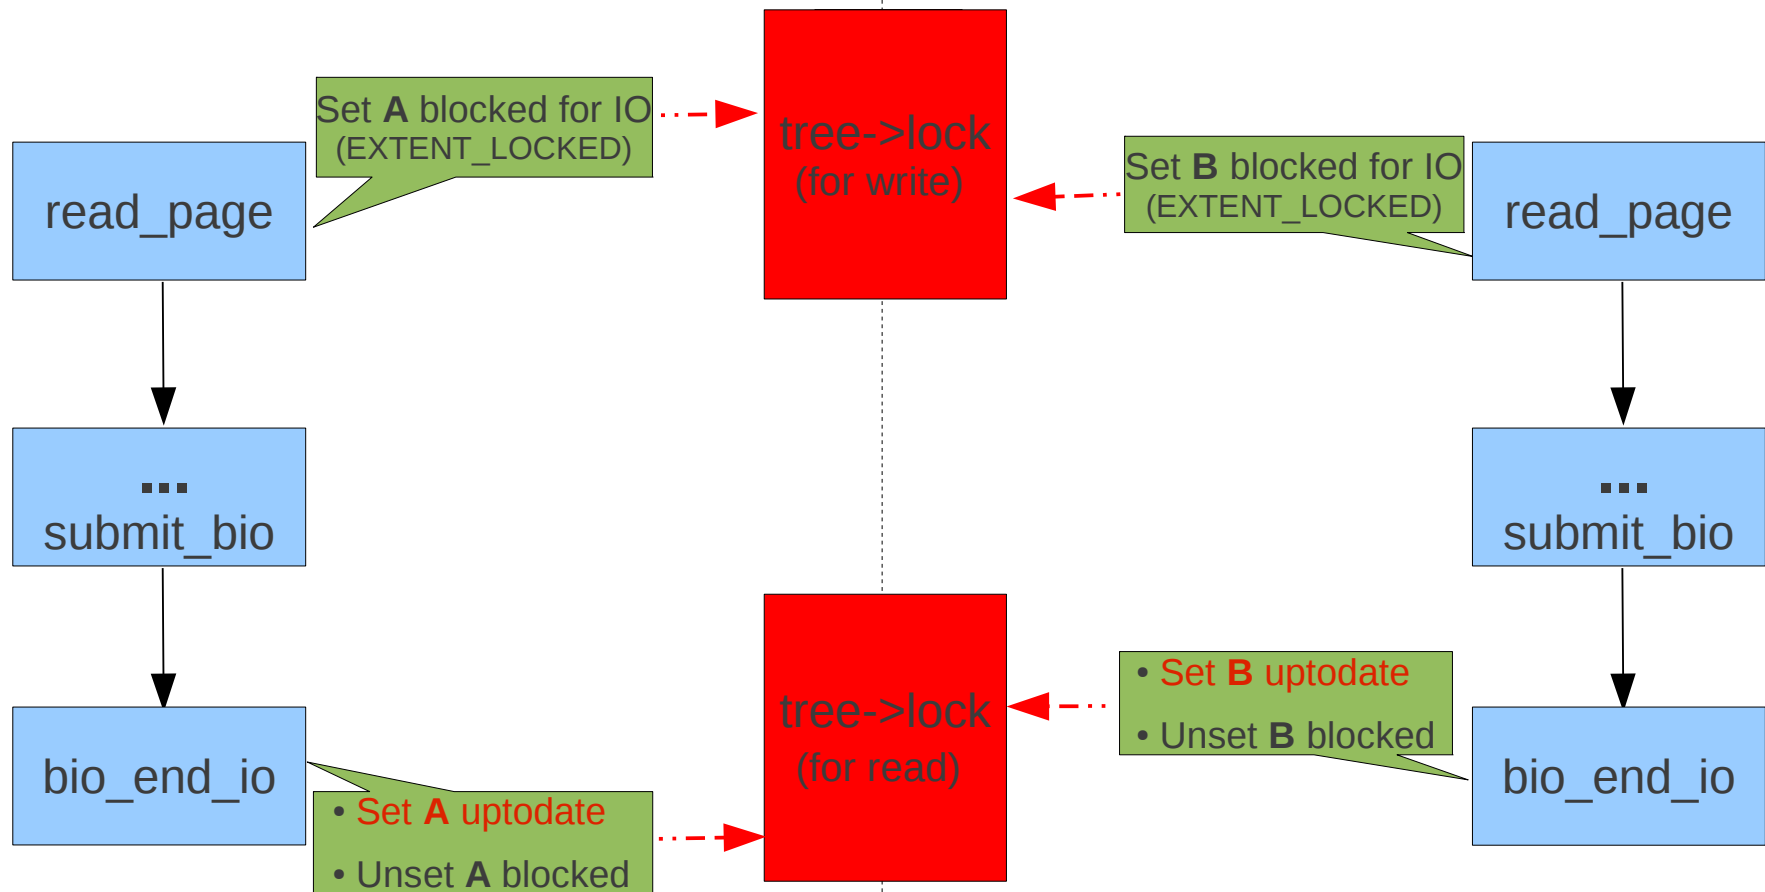
**Flags:**

- EXTENT\_DIRTY,
- EXTENT\_LOCKED,
- EXTENT\_UPTODATE,
- EXTENT\_DEALLOC,
- etc.

Extent state tree rbrroot



Race on **tree's lock** between **A** and **B**



# Problems we're facing with


- Over reservation leads to ENOSPC
- Lock contention in kernel data structures
  - Resort to rcu + rbtree / btree / skiplist for lockless read?
  - Reduce lock granularity

# Problems we're facing with:

Over reservation leads to ENOSPC

- Btrfs is based on B+ tree
- COW on WAFL
- We are not able to know the precise number of space we're going to use

**Because...**



**Over  
reservation**



**Then...we have some available space that cannot be allocated :(**

# Problems we're facing with:

## Lock contention in kernel data structures

- Lock contention in in-core rbtrees
  - Extent state tree
  - Free space tree
  - etc
- Possible ways for lockless read
  - Probabilistic skiplist with RCU lock
  - Rbtree with RCU lock
  - Btree with RCU lock
  - Smaller lock granularity

# Future work

- Fork a buddy system on space allocation
- Lockless metadata
- Btrfsck (offline/online)
- Performance
  - overall better than ext3 and ext4

# The whys and wherefores of using btrfs

- Good performance
- Good scalability
- Good reliability
- High fault tolerance
- Ease of management
- Base stone of distributed file systems like Ceph, etc.



# Thanks!

- Thank Liu Bo
- [btrfs.wiki.kernel.org](http://btrfs.wiki.kernel.org)