



Standards, the kernel and Open Source

<http://veillard.com/Talks/CLKBeijing2012.pdf>

Daniel Veillard
veillard@redhat.com

Standards: why ?

Main purpose is **interoperability**

- Public description of the technology
- Test suite or conformance checks

Main benefit is **cheaper, ubiquitous** technology

- Not tied to one vendor
- Larger user base
- Competition between vendors

Classic example: camera/phone cable



Standards: How ?

Various standard bodies, process is usually:

- 1) Multiple actors
- 2) Agreement to make a public specification
- 3) Shared work on creating that specification
- 4) Multiple initial versions (with feedback)
- 5) Vote and publication as a standard
- 6) Maintenance

Usually takes a few year, often painful

Standardization groups



And many others ...

What about the Linux kernel ?

- Linux is in the C language, which is standardized !
- Linux initial success was tied to the POSIX API
 - Implementing the standard gained a lot of applications
- We rely on standardized hardware
 - Buses (PCI/I2C...)
 - Protocols to talk to disks and other peripherals
 - Boot process
- We rely on standardized networking
 - From the lowest level: frame/packet level
 - Up to the application: Web, Video, etc ...

Looking more closely

Looking at the linux kernel code 3.6.1 source code


- IETF RFC standards:
 - Reference 212 different RFCs in the code base (544, 791, 792, 793, ... up to 5961, 6106, 6164, 6298)
- Many many references to IEEE
 - 802.11/802.15.4 for all the wireless
 - 1394/1212/1284 for firewire/SCSI/parallel
 - 754 floating point arithmetic
- ISO standards:
 - CD filesystems, character sets, networking
- ...

Why should I care ?

- Sometimes you won't have the choice:
 - Interoperability is crucial
 - New hardware is coming
- Sometimes you have the choice:
 - Which standard(s) to implement
 - Finding relevant standards can be challenging
 - You may not like it, make sure you don't exclude it by design
- Sometimes you want to be involved:
 - It may still be time to fix it !
 - They may need your implementation to finish
 - Providing test case and suites helps interop

Standard and Software in parallel


- Parallel developments
 - Same deadline
 - Reference code
- Good points:
 - Feedback
 - Timing is good
 - Positive perception
- Bad point:
 - What if the standard doesn't pass
 - Frequent changes to the code as the draft evolves



That situation is not very common

Late implementor

- React upon demand
 - Existing need
 - Spec looks okay
 - Good points:
 - Spec is stable
 - Existing User base
 - Minimal effort
 - Benefit from earlier implementor efforts
 - Bad points:
 - Too late to change the specification
 - Competing with existing implementations
- That situation is very common, usually the easiest



Early implementor

- Standardize existing code base(s)

- Software works
- Build a standard

- Good points:


- Clear direction
- Existing User base

- Bad points:

Existing user base

- The specification will change, your code too
- Competing with other people on the standard choices

That happens, this can be hell



Open Source specific

- Our code is public, are the standards (or drafts) too ?
- Do we have the resources to implement the spec fully ?
- Collaboration with the Working Group can be great:
 - Feedback loop integrate the Open Source Process
 - Who pays for the membership fees ?
 - Can be very time consuming
- One very hard issue : **Patents**
 - Affects us harder than proprietary code
 - Different standard bodies approaches
 - Royalties free (W3C)
 - RAND (Reasonable non discriminatory)
 - Workarounds are not always possible

Conclusions

- A lot of standards impact Open Source projects
 - Don't ignore them, be ready
 - Sometimes it is worth contributing
 - Be careful in your implementation
 - Interop is important
 - Avoid legal issues
- Some standard body are friendly to OSS
 - Free 'expert' access
 - Legal provisions to avoid patent issues

<http://veillard.com/Talks/CLKBeijing2012.pdf>