

王贊 from China LTC

Michael Wang <wangyun@linux.vnet.ibm.com>



Smart Wake-Affine In Scheduler

Agenda

Terminology

Background

Benefit & Issue & Analysis

Figure out the Solution

Our solution in mainline

Thinking

Terminology

wakeup

un-running process need to be waken up to run

waker

The running process which try to wakeup an un-running process

wakee

un-running process to be wakeup

neighbor CPU

Logical cpus which belong to the same core or the same socket

running closely

Processes run on neighbor CPU

related processes

Processes access the same piece of physical memory

pull

Choose a neighbor CPU of waker's to run wakee

Background

What is wake-affine?

Feature inside scheduler, by which we attempt to make processes running closely

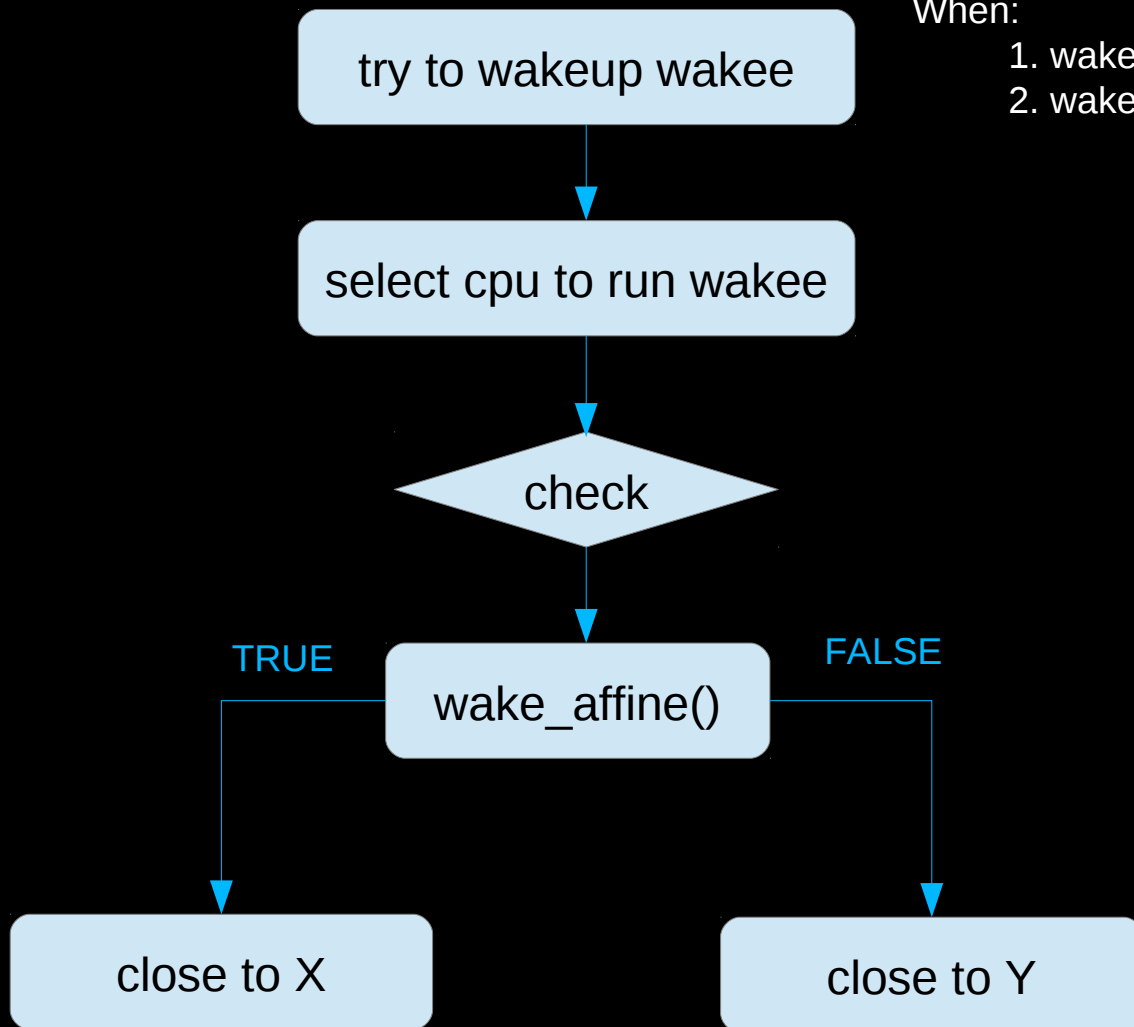
Why wake-affine?

In some cases, related processes running closely will gain benefit, mostly from cache-hit

How wake-affine works?

When:

1. waker is currently running on CPU X
2. wakee was last time running on CPU Y

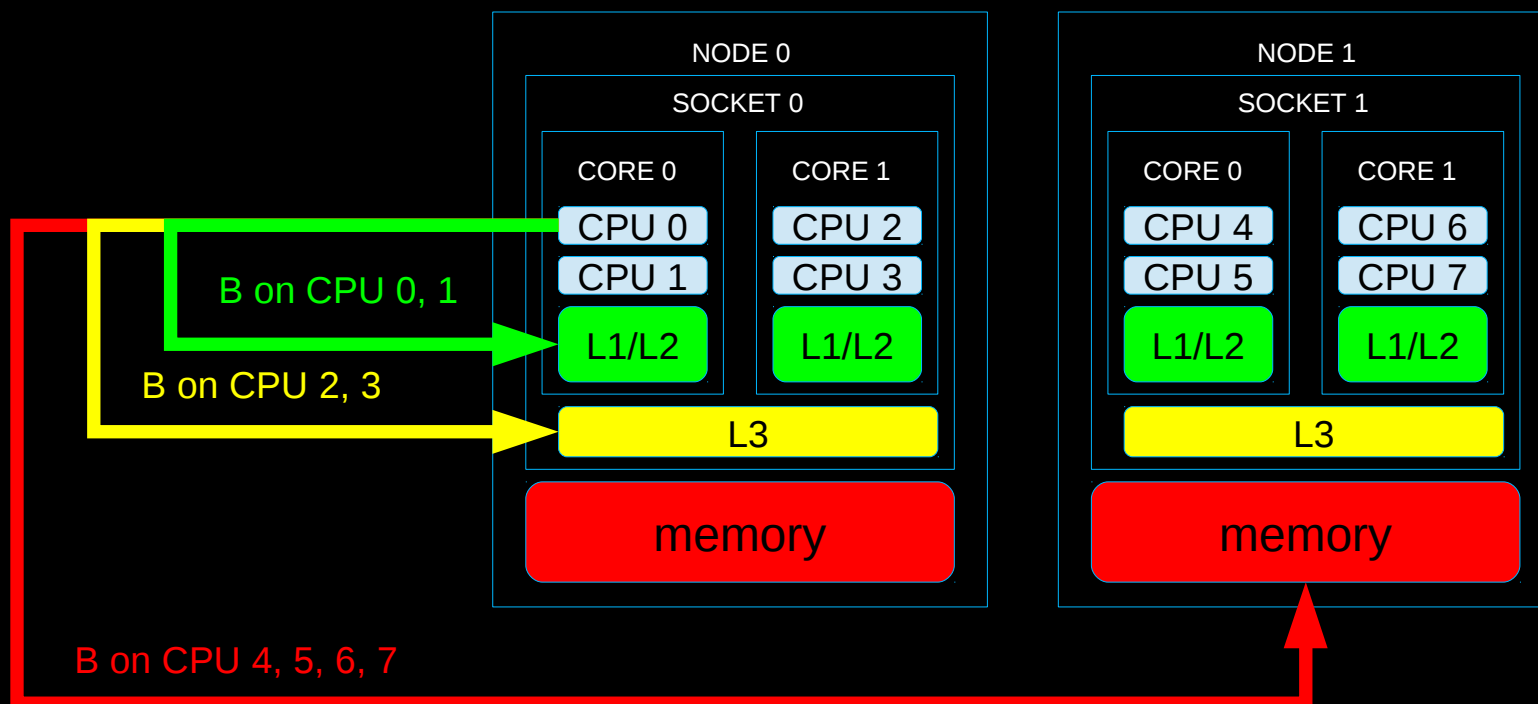


How distance influence the performance?

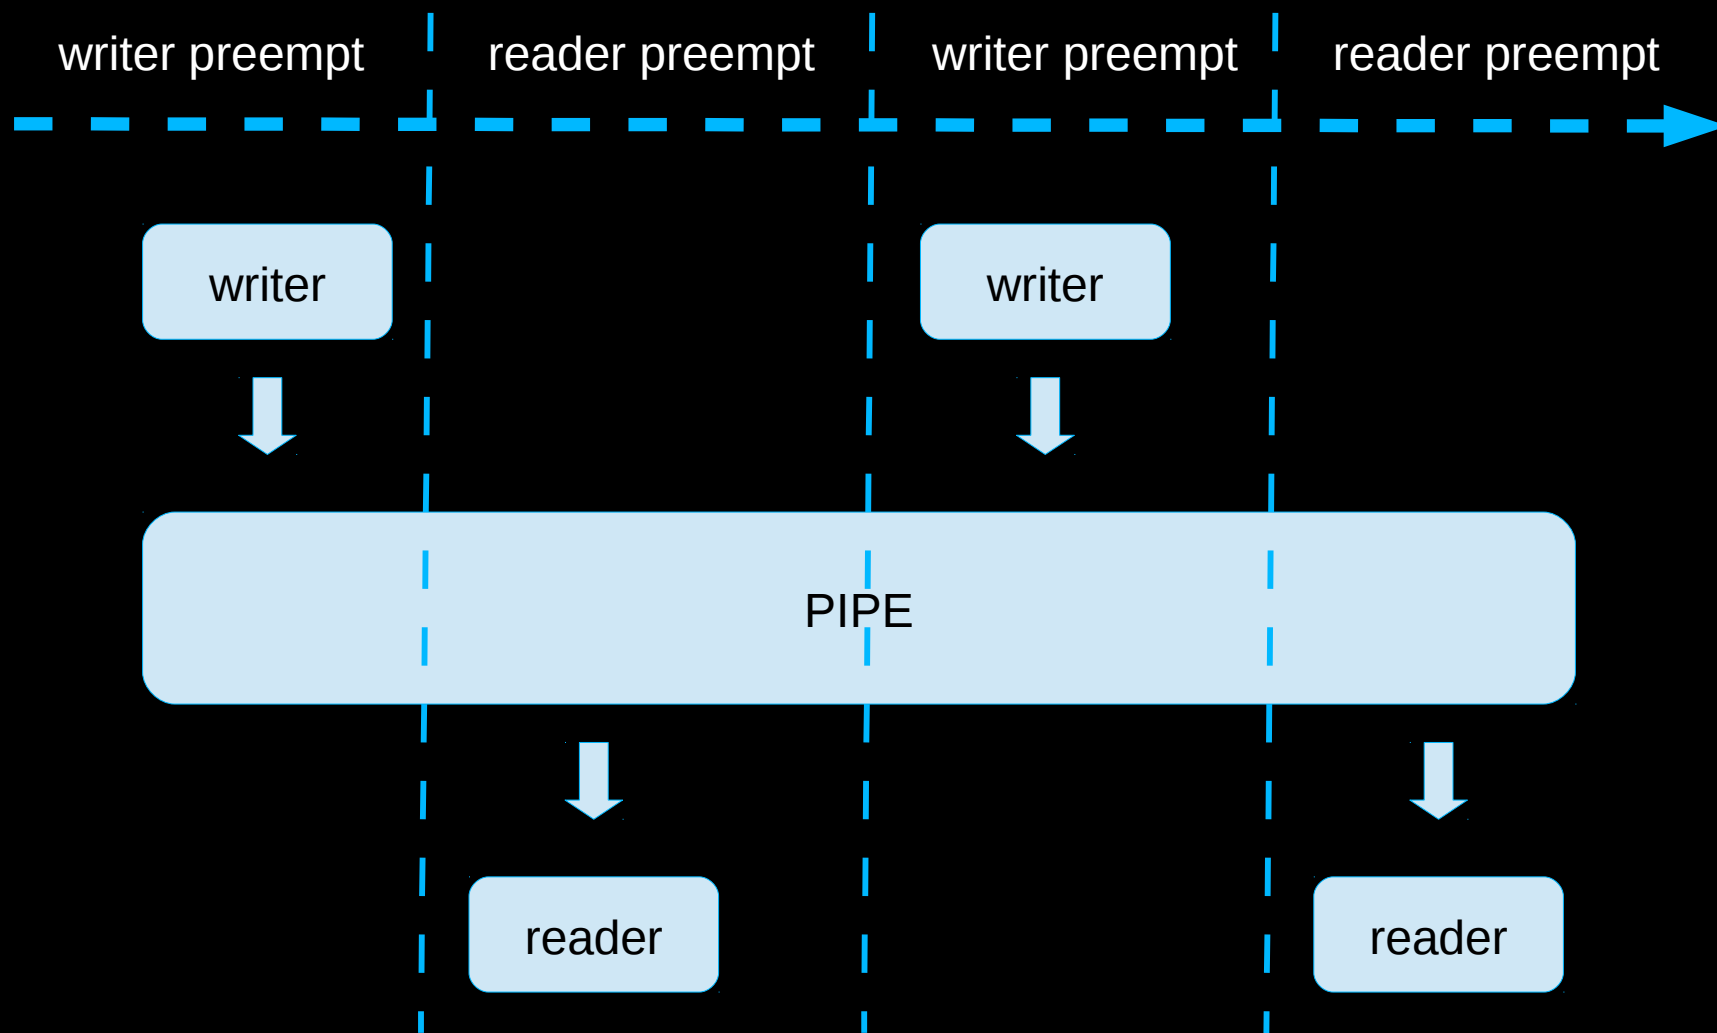
When:

1. process A is running on CPU 0
2. process B accessed memory of NODE 1
3. the accessed memory has been cached
4. process A want to access the same memory

Then when B run on different CPU...



Special case when writer and reader on same CPU



No wakeup needed

Benefit & Issue & analysis

Benefit

during testing, we found that this feature improved hackbench around
15%

Issue


during testing, we found that this feature decrease pgbench around
40%

Analysis

we noticed that pgbench is a benchmark for database, which means the request from clients may finally come into one server, their wakeup relationship are different

wake up relationship difference

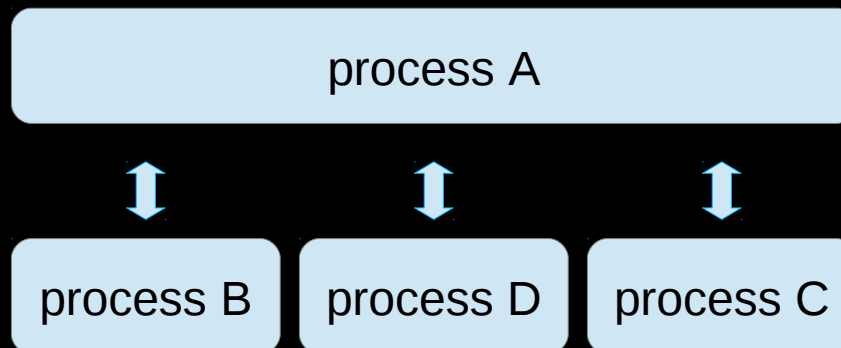
N : N



A, D will wakeup or be waken by B, C.

A pull B, C
B pull A, D
C pull A, D
D pull B, C

1 : N



A will wakeup and be waken by others.

A pull B, C, D
B pull A
C pull A
D pull A


The bad influence of wake-affine

CPU of NODE 0


CPU of NODE 1

N : N

After writer x wakeup readers



After writer y wakeup readers



1 : N


After server wakeup x



After server wakeup y



After server wakeup z



Analysis

So the issue is that wake-affine has the potential to cause starvation in 1 : N cases, the lost performance caused by:

- 1. overhead of wake-affine**
- 2. wrong decision of wake-affine which make server starving**
- 3. chain reaction caused by server which delay all it's clients**

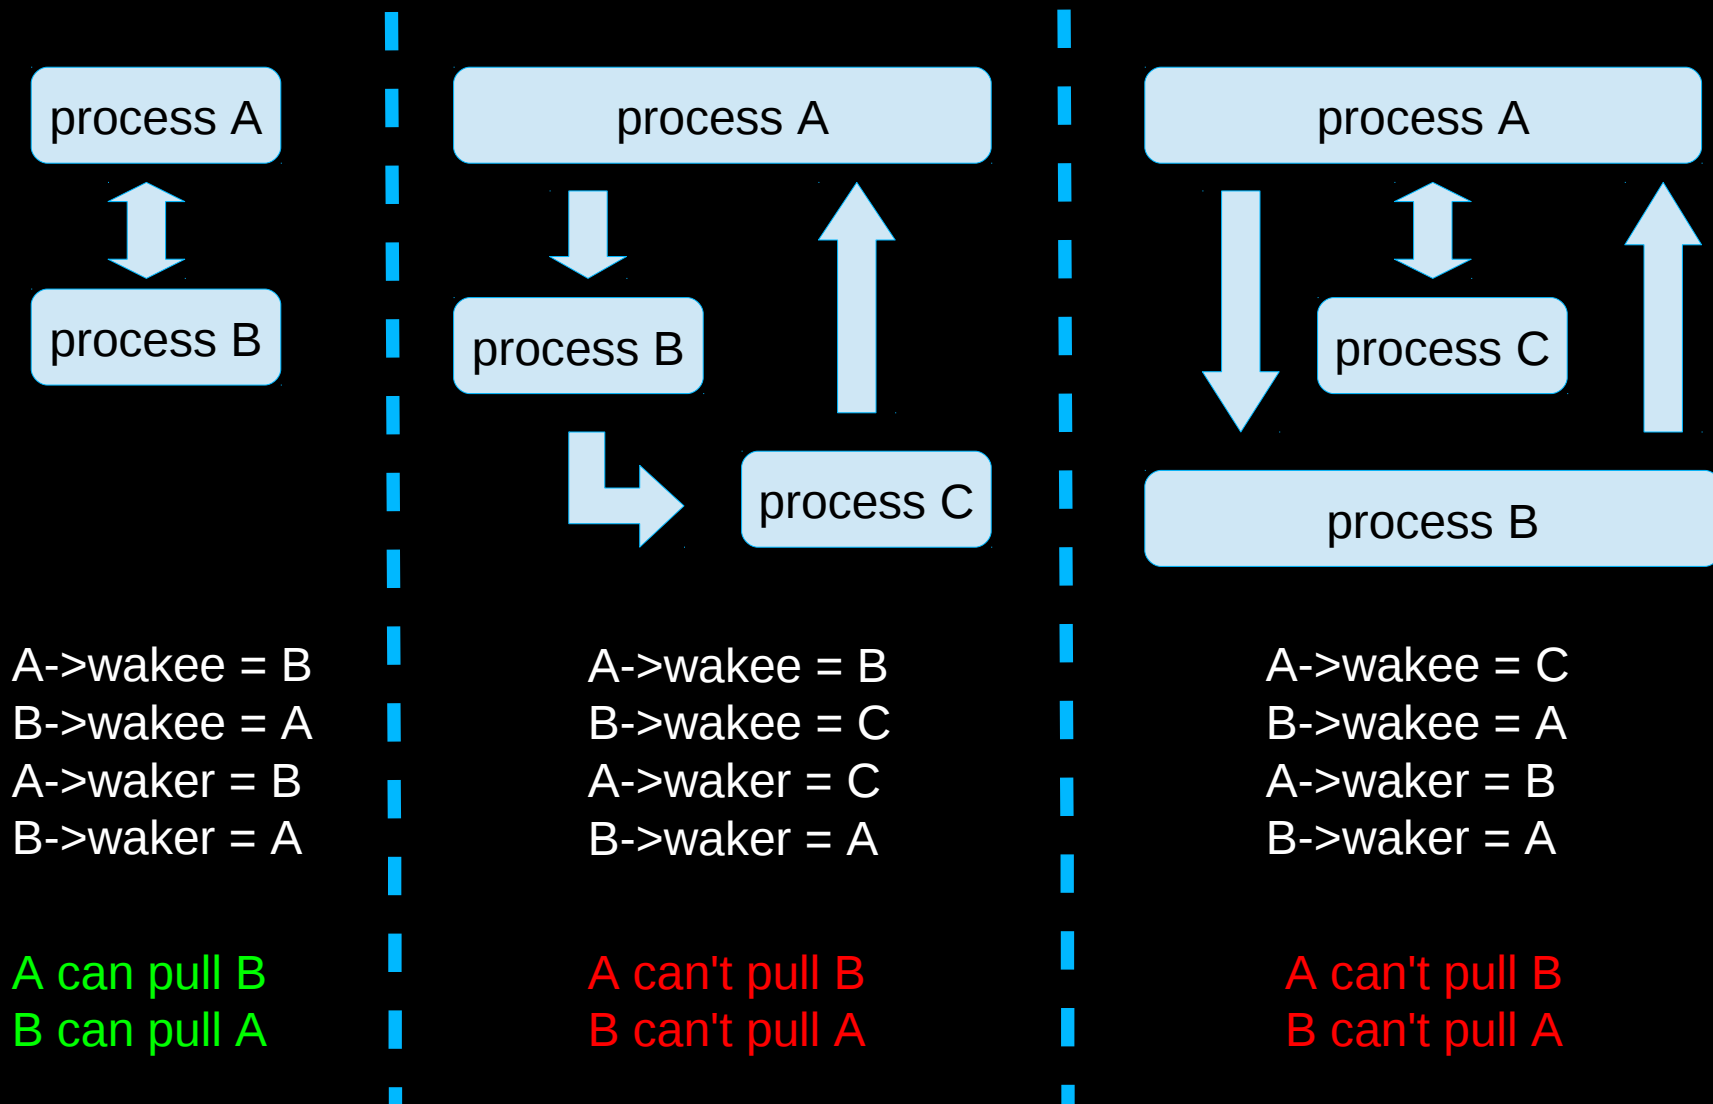
Figure out the solution

Remove wake-affine logical
very easy to implement
hackbench will lost it's benefit

Throttle wake-affine
easy to implement
hard for user to use

Buddy wake-affine
only allow to pull when waker and wakee are keeping wakeup each other.
accurately address the related processes

Buddy wake-affine



Buddy wake-affine

Still facing two issues:

- 1. overhead is too high.**
- 2. lost target in soft-irq context.**

Thus such a solution is required:

- 1. overhead very low.**
- 2. automatically detect the bad cases, meanwhile won't kill the good cases.**


Our solution in mainline

Smart wake-affine

This is a variant of buddy wake-affine solution, recording the flip of wakee switching, stop wake-affine when waker has high flip.

Flip difference

N : N



one round flips

A = 2

B = 2

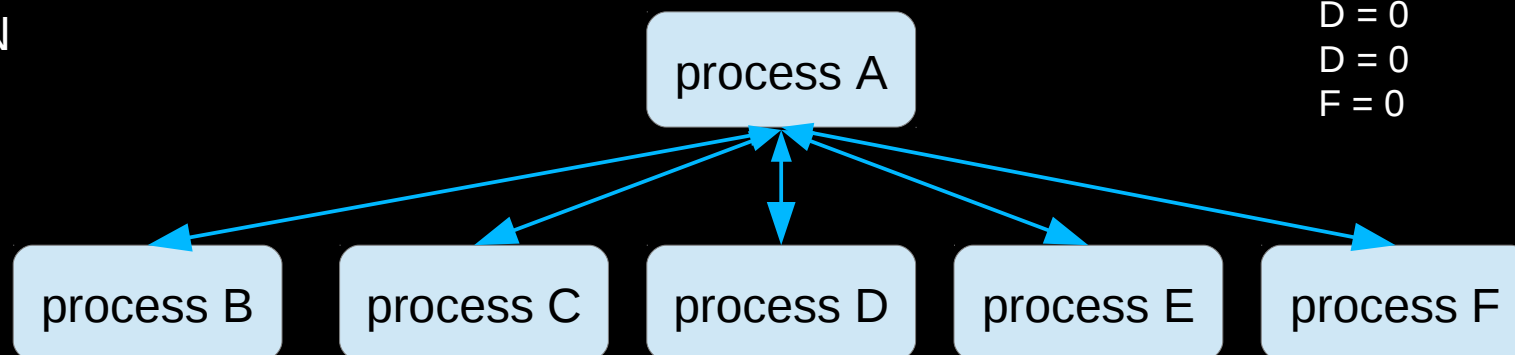
C = 2

D = 2

E = 2

F = 2

1 : N



one round flips

A = 4

B = 0

C = 0

D = 0

E = 0

F = 0

Our solution in mainline

```
static void record_wakee(struct task_struct *p)
{
    if (jiffies > current->wakee_flip_decay_ts + HZ) {
        current->wakee_flips = 0;
        current->wakee_flip_decay_ts = jiffies;
    }

    if (current->last_wakee != p) {
        current->last_wakee = p;
        current->wakee_flips++;
    }
}
```

```
static int wake_wide(struct task_struct *p)
{
    int factor = this_cpu_read(sd_llc_size); //top cache-share domain size, added by PeterZ for this solution
    if (p->wakee_flips > factor)
        if (current->wakee_flips > (factor * p->wakee_flips)) //bigger NODE can hold more threads
            return 1;
    return 0;
}
```

Our solution in mainline

Results:

- overhead is low**
- automatically adjust the scope**
- get back the lost 40% performance of pgbench**
- reserve the 15% gotten performance of hackbench**
- netperf show 3-5% increased performance**

Thanks to Fengguang Wu for the data of netperf :)

Thinking

Scheduler used to rely on load-info to make sure it's fairness, but this case show us that only take care the load is not enough in some situation, like this case, process with same load but different wakeup flips should be handled differently.

There are still a wide variety of characteristic we may could use to drive the scheduler decision, by which make it more smart while handling some situation and without conflict with the others.

By analyze the workload of APP, we may also be able to address the key point for optimization, usually the bottleneck and the reason to make it the bottleneck.

Links

commit 7d9ffa8961482232d964173cccba6e14d2d543b2

Author: Peter Zijlstra <peterz@infradead.org>

sched: Micro-optimize the smart wake-affine logic

<https://lkml.org/lkml/2013/7/23/866>

commit 62470419e993f8d9d93db0effd3af4296ecb79a5

Author: Michael Wang <>wangyun@linux.vnet.ibm.com>

sched: Implement smarter wake-affine logic

<https://lkml.org/lkml/2013/7/23/875>

Thanks to all the folks who involved, special thanks to Mike Galbraith and Peter Zijlstra.

Thanks ;-)

王贊 from China LTC

Michael Wang <wangyun@linux.vnet.ibm.com>

© 2013 IBM Corporation