



中国Linux内核开发者大会

2016

Linux Performance Profiling & Visualization

Barry Song & Bob Liu & Mac Xu

performance measurements



top

iostat

vmstat

Sar oprofile

perf

...

Problems in existing tools

RAW data, difficult for non-experts to understand

e.g. load average...

```
top - 18:10:33 up 12 min,  2 users,  load average: 0.24, 0.40, 0.39
Tasks: 204 total,   1 running, 203 sleeping,   0 stopped,   0 zombie
%Cpu0  :  8.9 us,  1.0 sy,  0.0 ni, 90.1 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu1  :  5.1 us,  0.3 sy,  0.0 ni, 94.6 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
Mem:    1070116 total,   661600 used,   408516 free,   171560 buffers
```

Problems in existing tools

Not visualized

e.g. how many memory used in Linux?

```
baohua@baohua-VirtualBox:~$ free
```

	total	used	free	shared	buffers	cached
Mem:	1079116	863188	215928	3068	171600	305600
-/+ buffers/cache:		385988	693128			
Swap:	522236	0	522236			

Problems in existing tools

Lack of the description for changes

e.g. how the CPU usage is changing during a period?

```
top - 18:13:55 up 16 min,  2 users,  load average: 0.05, 0.23, 0.32
Tasks: 204 total,  1 running, 203 sleeping,  0 stopped,  0 zombie
%Cpu(s):  4.6 us,  0.3 sy,  0.0 ni, 95.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem: 1079116 total,  869656 used,  209460 free,  171608 buffers
KiB Swap:  522236 total,    0 used,  522236 free.  305616 cached Mem
```


Problems in existing tools

Cannot interact with users smoothly

e.g. what if we only care about a particular process in "top" ?

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3005	baohua	20	0	157392	79036	35196	S	6.0	7.3	0:20.68	compiz
1979	root	20	0	115640	49936	16792	S	2.3	4.6	0:08.01	Xorg
1771	root	0	-20	3844	3828	2636	S	1.3	0.4	0:00.17	atop
3425	baohua	20	0	129308	32356	24324	S	1.0	3.0	0:01.56	gnome-ter+
3595	baohua	20	0	7012	3176	2768	R	0.7	0.3	0:00.11	top
3045	baohua	20	0	38300	7808	5280	S	0.3	0.7	0:00.10	indicator+
1	root	20	0	4632	3752	2556	S	0.0	0.3	0:02.52	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.14	ksoftirqd+
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0+
6	root	20	0	0	0	0	S	0.0	0.0	0:00.10	kworker/u+
7	root	20	0	0	0	0	S	0.0	0.0	0:00.38	rcu_sched
8	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_bh
9	root	rt	0	0	0	0	S	0.0	0.0	0:00.02	migration+
10	root	rt	0	0	0	0	S	0.0	0.0	0:00.01	watchdog/0
11	root	rt	0	0	0	0	S	0.0	0.0	0:00.00	watchdog/1
12	root	rt	0	0	0	0	S	0.0	0.0	0:00.02	migration+
13	root	20	0	0	0	0	S	0.0	0.0	0:00.08	ksoftirqd+

A Solution

- Intuitive
- Visualized
- Interactive
- Ease of Use

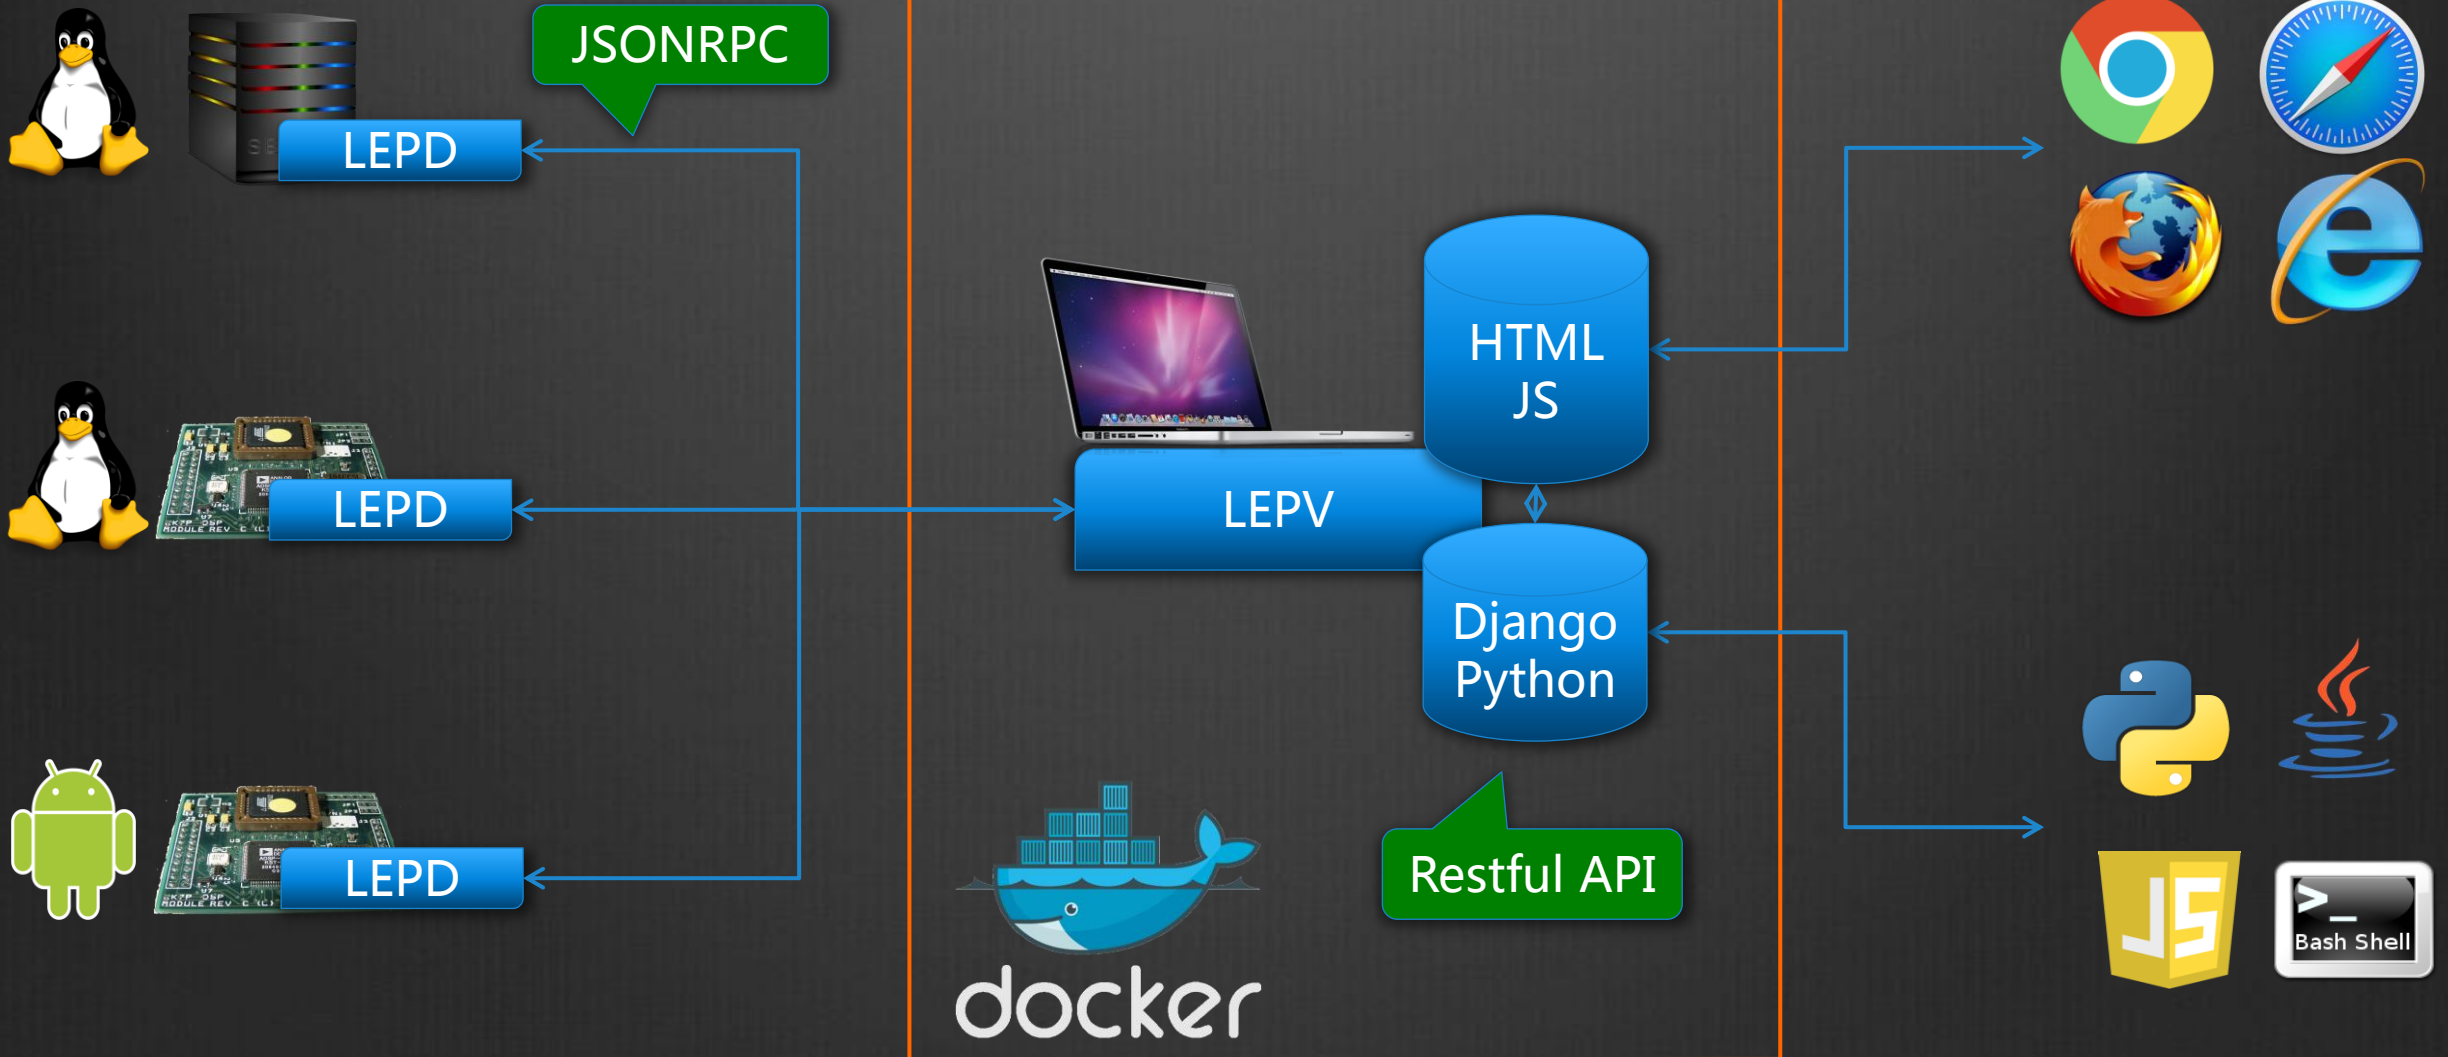
LEP

What is LEP

Linux Easy Profiling

- Web-based
- Open-source
- All-In-One

LEP architecture



LEP Summary

- C/S
- PC / Embedded Board
- Linux / Android
- JSONRPC
- Web App
- Restful API
- Docker

LEP Dataflow


HTTP GET
AVG LOAD



django



JSON RPC
AVG LOAD



```
{  
  'last15':Decimal('2.13'),  
  'last5':Decimal('2.25'),  
  'last1':Decimal('2.58')  
}
```

```
{  
  'result':'2.58 2.25 2.13 4/110 19674\n'  
}
```

/proc/loadavg

LEPV


LEPD

Live Demo




LEP: Load balance view


CPU Stat: Idle



CPU Stat: User+Sys+Nice




LEP: Memory consumption view




LEP: App memory usage view

应用程序耗费内存占总内存比例




应用程序内存消耗比例分布(基于PSS)



LEP: Average load view

Average Load



LEP: Symbol level view

基于Symbol的时间分布 (perf top)

11

Search: <input type="text"/>			
Command	Overhead	Shared Object	Symbol
malloc	53.53%	malloc	[.] main
swapper	32.87%	[kernel.kallsyms]	[k] native_safe_halt
malloc	12.50%	[kernel.kallsyms]	[k] clear_page
malloc	0.15%	[kernel.kallsyms]	[k] free_pages_prepare
malloc	0.13%	[kernel.kallsyms]	[k] get_page_from_freelist
malloc	0.10%	[kernel.kallsyms]	[k] _cond_resched
malloc	0.08%	[kernel.kallsyms]	[k] clear_huge_page
malloc	0.05%	[kernel.kallsyms]	[k] __do_page_fault
malloc	0.05%	[kernel.kallsyms]	[k] _raw_spin_lock
swapper	0.05%	[kernel.kallsyms]	[k] __do_softirq
malloc	0.03%	[kernel.kallsyms]	[k] page_add_new_anon_rmap
swapper	0.03%	[kernel.kallsyms]	[k] refresh_cpu_vm_stats
free	0.02%	[kernel.kallsyms]	[k] do_exit
free	0.02%	[kernel.kallsyms]	[k] page_add_file_rmap
free	0.02%	ld-2.19.so	[.] 0x00000000000009e2a
free	0.02%	libc-2.19.so	[.] 0x0000000000007eae8
gapd	0.02%	[kernel.kallsyms]	[k] __fdget_raw
gapd	0.02%	[kernel.kallsyms]	[k] FUTEX_WAIT
malloc	0.02%	[kernel.kallsyms]	[k] __do_softirq
malloc	0.02%	[kernel.kallsyms]	[k] lru_cache_add

LEP: Interaction with users

CPU TOP										⏏ ⚙
										Search: <input type="text" value="load"/>
PID	User	PRI	NI	VSZ	RSS	S	%CPU	▼ %MEM	Time	Command
4351	ubuntu	19	0	4196	624	S	38.2	0.0	00:14:28	./load
4352	ubuntu	19	0	4196	88	R	38.2	0.0	00:14:28	./load
4483	ubuntu	19	0	4196	88	S	38.2	0.0	00:14:20	./load
4482	ubuntu	19	0	4196	644	R	38.1	0.0	00:14:20	./load

LEP: Interaction with users

The screenshot displays the LEP (Linux Environment Profiler) interface. A modal dialog box titled "参数设置" (Parameter Settings) is centered on the screen. The dialog contains two input fields: "刷新频率 (单位: 秒)" (Refresh frequency (unit: seconds)) with a value of 5, and "最多显示数据数" (Maximum number of data items to display) with a value of 25. At the bottom right of the dialog are two buttons: "取消" (Cancel) in orange and "保存" (Save) in green. The background interface is dimmed, showing a "CPU TOP" header and a table of process data. A search bar is visible in the top right of the background interface. A yellow circle highlights a settings icon in the top right corner of the background interface.

参数设置

刷新频率 (单位: 秒) 5

最多显示数据数 25

取消 保存

CPU TOP

Search:

PID	User	PRI	NI	VSZ	RSS	S	%CPU	%MEM	Time	Command
4351	ubuntu	19	0	4196	624	R	38.2	0.0	00:15:59	./load
4352	ubuntu	19	0	4196	88	S	38.2	0.0	00:16:00	./load
4482	ubuntu	19	0	4196	644	R	38.2	0.0	00:15:52	./load

Next steps for LEP

- More features
- Better documentations

Upcoming features

- ARM support
- Predict memory leak for an application and kernel
- Analyze cache miss
- Analyze time consumption for one native/Java processes
- Benchmark integration
- I/O queues
- Kernel memory details such as buddy, slab, CMA etc.
- Runtime scheduler(CPU/IO)
- Boot procedure
-

Plans - LE series

LEB

Easy Building

LET

Easy Testing

LEP

Easy Profiling

Contributors

- Barry Song
- Bob Liu
- Mac Xu
- More developers are coming
 - Ping Liu
 - Ray Chen
 - Joy Zhou

➤ YOU

Fork the code here:

- LEPD:
<git@www.linuxep.com:repo/lep/lepd>
- LEPV:
<git@www.linuxep.com:repo/lep/lepv>



Q&A