



# arm

## GIC-based Interrupt handling mechanism on AArch64

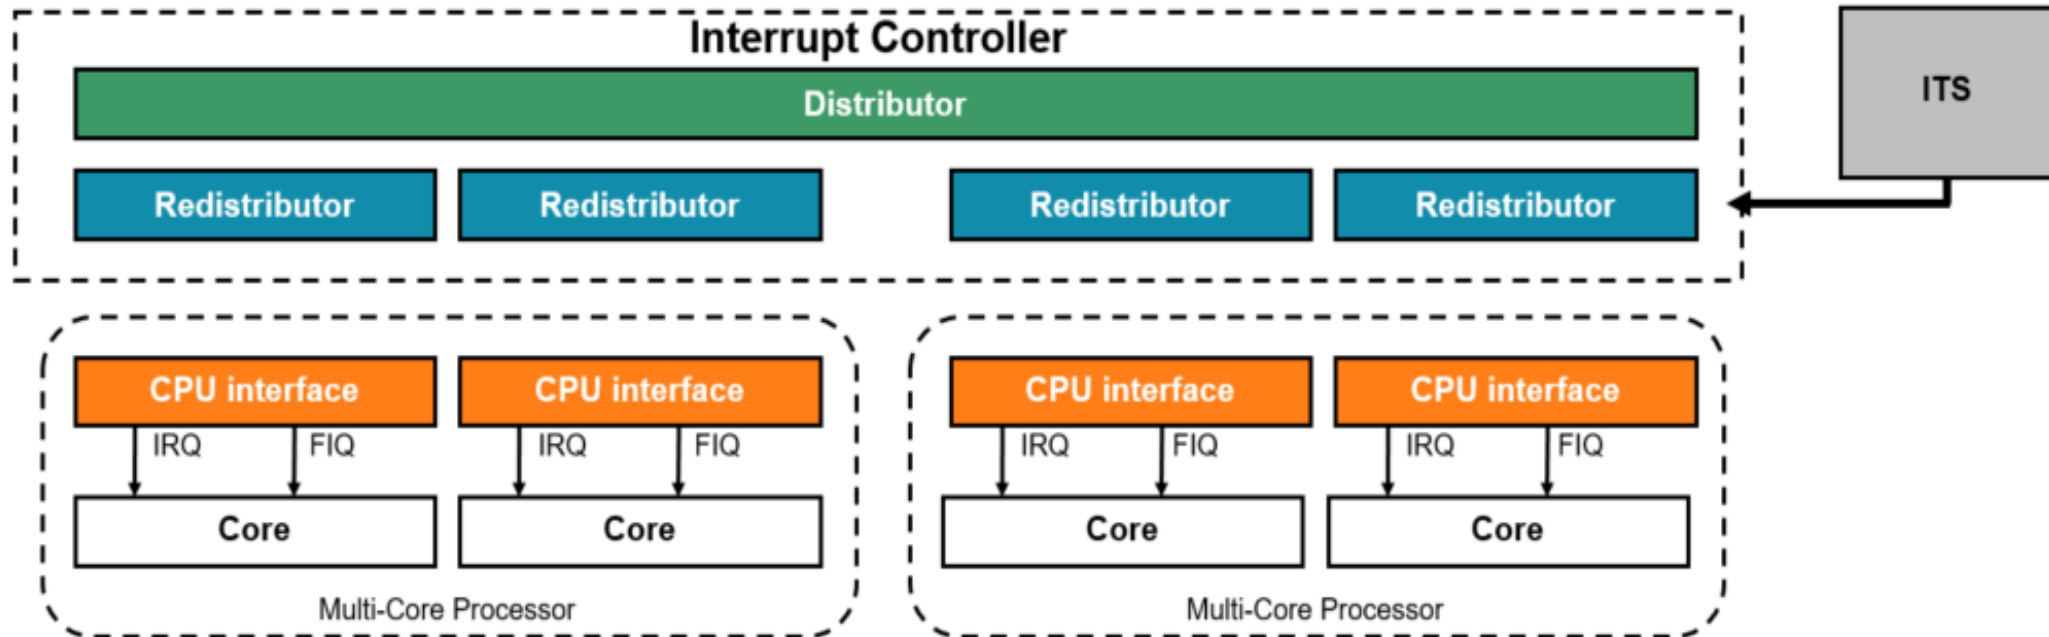
Dennis Chen @ Arm

# Agenda

- Part 1 - GICv3 Introduction
- Part 2 - Interrupt handling on AArch64
- Q & A

# GLCv3 Architecture

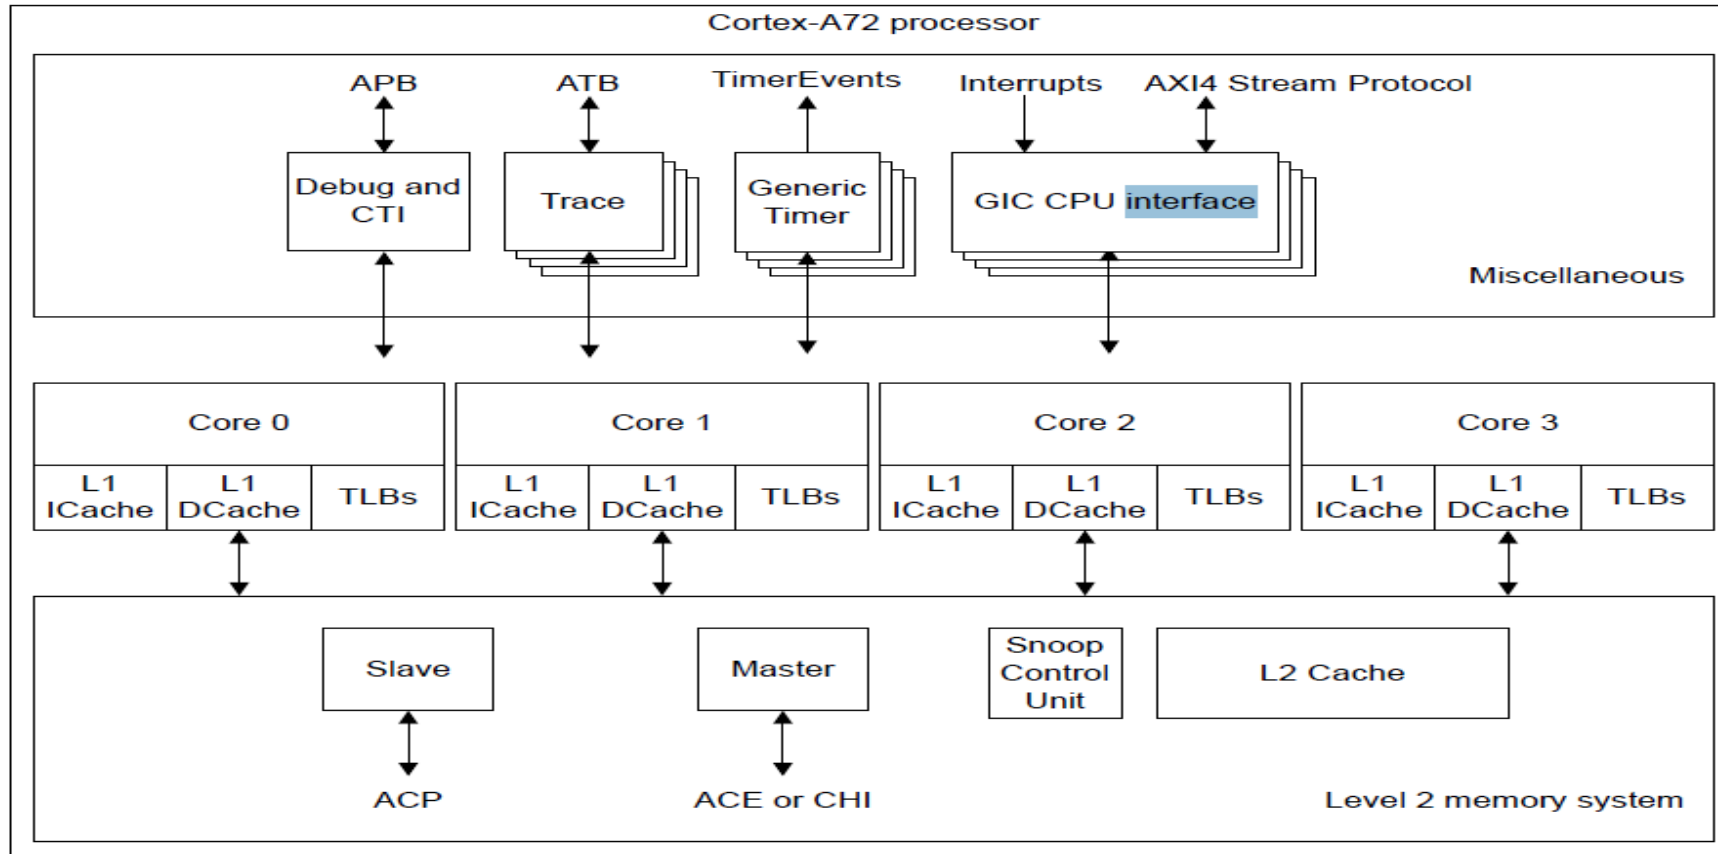
- The Generic Interrupt Controller (GIC) Architecture defines a programmers model for interrupt controllers



- Distributor and Redistributor interfaces are memory mapped
- CPU interface accessed as system registers

# GICv3 Architecture

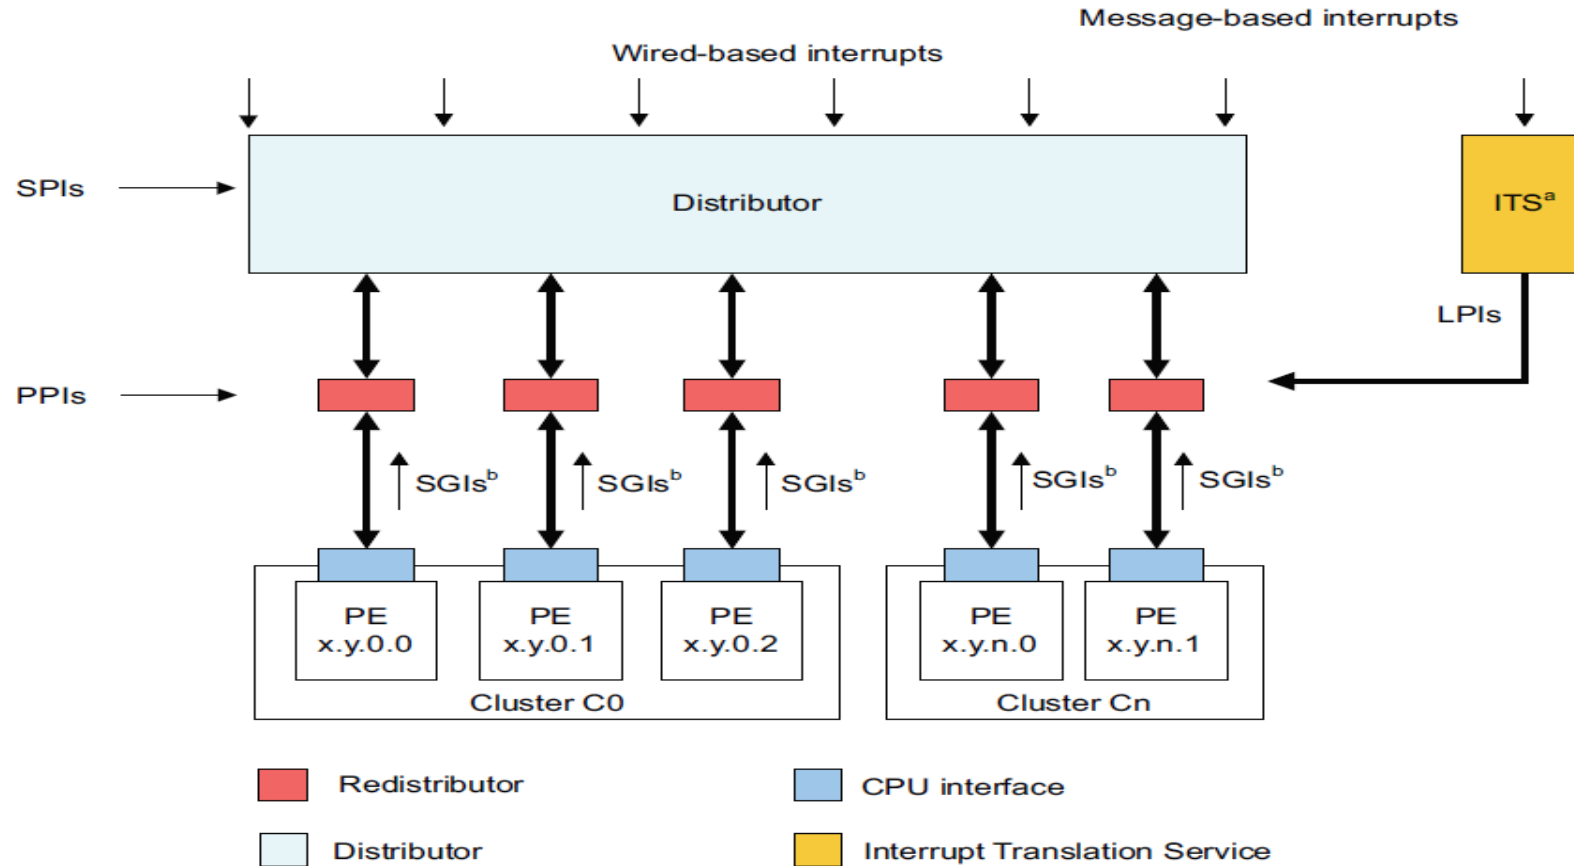
- The GIC CPU interface in an Arm MPCore Processor



# GIC interrupt types

- **SPI – Shared Peripheral Interrupt**
  - Peripheral interrupt, available to all the cores using the interrupt controller
  - INTIDs 32 to 1019
- **PPI – Private Peripheral Interrupt**
  - Peripheral interrupt which is private to an individual core
  - INTIDs 16 to 31
- **LPI – Locality-specific Peripheral Interrupt** *(new in GICv3)*
  - Peripheral interrupt, typically routed by an ITS
  - INTIDs 8192+
- **SGI – Software Generated Interrupt**
  - Triggered by writing to a register within the interrupt controller
  - INTIDs 0 to 15
- **INTIDs 1020 to 1023 are reserved and have special purposes**
  - Take 1023 as an example. INTID=1023 means no pending interrupt.

# GIC logic partition




a. The inclusion of an ITS is optional, and there might be more than one ITS in a GIC.

b. SGIs are generated by a PE and routed through the Distributor.

# Message based interrupt – new in GICv3

- **GICv3 adds support for message based interrupts (MBI)**
  - Instead of using a dedicated signal, a peripheral writes a register in the GIC to trigger an interrupt



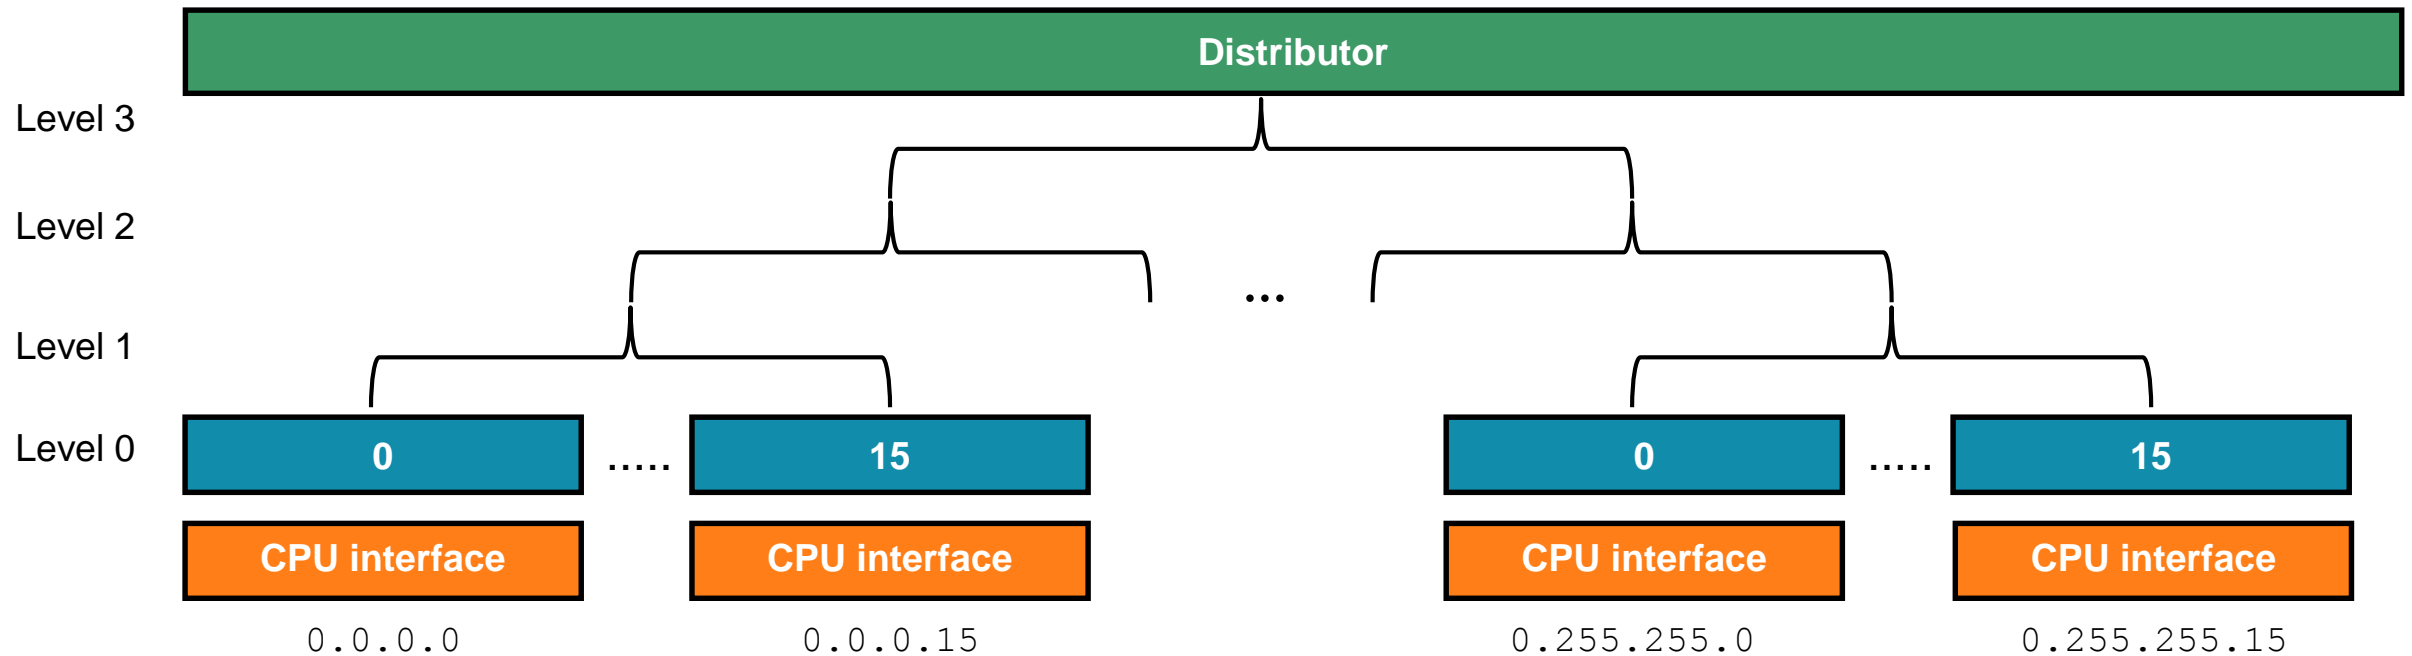
- **Why?**
  - Can reduce the number of wires needed and ease routing
  - Matches model used by PCIe
- **In most cases software should not care whether interrupt is an MBI or not**
  - Does not change how state machine operates

# Interrupt configuration

- **Each INTID has a number of settings associated with it:**
- **Enable**
  - Only enabled interrupts can be forwarded to a core
  - However, disabled interrupts can still become pending if the source is asserted
- **Priority**
  - Each interrupt has an 8-bit priority associated with it
  - 0x00 is the highest priority, 0xFF is the lowest priority
  - A GIC can implement less than 8 bits of priority, if so only the most significant bits are used
- **Configuration**
  - Whether the interrupt is level-sensitive or edge-triggered
- **Security/Group and Target**
  - Covered on the next few slides



# SPI routing in GICv3 – Affinity levels



- **GICv3 identifies the attached CPU using a four level ID, called affinity**
  - Matches affinity mapping defined for ARMv8-A, reported in the `MPIDR_EL1`
- **For each SPI there is a `GICD_IROUTERn` register which controls where the interrupt is routed to**
  - Two options:
    - Send to any connected CPU (referred to as "1 of N")
    - Send to one specific CPU (affinity coordinates of CPU written to register)

# Interrupt states

- **Inactive**

- interrupt is not active and not pending

- **Pending**

- interrupt is asserted but not yet being serviced

- **Active**


- interrupt is being serviced but not yet complete

- **Active & Pending**

- interrupt is both active and pending

- **Interrupt goes:**

- Inactive → Pending      when the interrupt is asserted
  - Pending → Active      when a CPU acknowledges the interrupt by reading the Interrupt Acknowledge Register (IAR)
  - Active → Inactive      when the same CPU deactivates the interrupt by writing the End of Interrupt Register (EOIR)




a. Not applicable for LPIs.

# Interrupt Acknowledging

- **On taking an interrupt software must read one of the [Interrupt Acknowledge registers](#)**
  - This returns the INTID of the interrupt and updates the state machine
- **There are different registers for Group 0 and Group 1 interrupts**
  - [ICC\\_IAR0\\_EL1](#) – for Group 0 interrupts
  - [ICC\\_IAR1\\_EL1](#) – for Group 1 interrupts
- **When an interrupt is acknowledged the [Running Priority](#) of the CPU interface takes on the priority of the interrupt**
  - Current value reported by [ICC\\_RPR\\_EL1](#)
  - Matters for pre-emption, covered later

# ITS – PCI MSI/MSIx



- An Interrupt Translation Service (or ITS) maps interrupts to INTIDs and Redistributors
- How is an interrupt translated?
  - Peripheral sends interrupt as a message to the ITS
    - The message specifies the **DeviceID** (which peripheral) and an **EventID** (which interrupt from that peripheral)
  - ITS uses the DeviceID to index into the Device Table
    - Returns pointer to a peripheral specific **Interrupt Translation Table**
  - ITS uses the EventID to index into the Interrupt Translation Table
    - Returns the **INTID** and **Collection ID**
  - ITS uses the Collection ID to index into the Collection Table
    - Returns the **target Redistributor**
  - ITS forwards interrupt to Redistributor

# Interrupt entry on AArch64

- **Vector Table** – A number of consecutive word-aligned addresses in memory, starting at the *Vector Base Address*. Each EL has an associated *Vector Base Address Register (VBAR)*

Table D1-7 Vector offsets from vector table base address

Exception taken from	Offset for exception type			
	Synchronous	IRQ or vIRQ	FIQ or vFIQ	SError or vSError
Current Exception level with SP_EL0.	0x000	0x080	0x100	0x180
Current Exception level with SP_ELx, x>0.	0x200	0x280	0x300	0x380
Lower Exception level, where the implemented level immediately lower than the target level is using AArch64. <sup>a</sup>	0x400	0x480	0x500	0x580
Lower Exception level, where the implemented level immediately lower than the target level is using AArch32. <sup>a</sup>	0x600	0x680	0x700	0x780

- a. For exceptions taken to EL3, if EL2 is implemented, the level immediately lower than the target level is EL2 if the exception was taken from Non-secure state, but EL1 if the exception was taken from Secure EL1 or EL0.

*EL0 has no corresponding VBAR*

# Interrupt entry on AArch64

- **Vector Table establishment in kernel**

arch/arm64/kernel/entry.S:

```
.align      11
ENTRY(vectors)
...
ventry     el1_irq
...
END(vectors)
```

ventry el1\_irq is:

```
.align 7
b     el1_irq
```

The IRQ handling chain:

el1\_irq → irq\_handler → handle\_arch\_irq

drivers/irqchip/irq-gic-v3.c:

set\_handle\_irq(gic\_handle\_irq), so – handle\_arch\_irq = gic\_handle\_irq

The main interrupt handling entry function is:

```
static asmlinkage void __exception_irq_entry gic_handle_irq(struct pt_regs *regs)
{
    u32 irqnr; //INTID
    do {
        irqnr = gic_read_iar(); // Read ICC_IAR1_EL1, get INTID & ACK interrupt
        gic_write_eoir(irqnr);
        handle_irq();
    } while (irqnr != ICC_IAR1_EL1_SPURIOUS); // ICC_IAR1_EL1_SPURIOUS = 1023
}
```

***1023 is the special INTID indicating no pending interrupt.***

# MADT of GIC

- ACPI MADT (Multiple APIC Description Table) – Firmware Description of GIC topology/info

## MADT

Header	Signature 'APIC'
	. . .
	Creator Revision
Local Interrupt Controller Address	
Flags	
Interrupt Controller Structure[n]	0xB GIC CPU Interface(GICC)
	0xC GIC Distributor(GICD)
	0xD GIC MSI Frame
	0xE GIC Redistributor (GICR)
	0xF GIC Interrupt Translation Service (ITS)

Table 5-63 GICD Structure

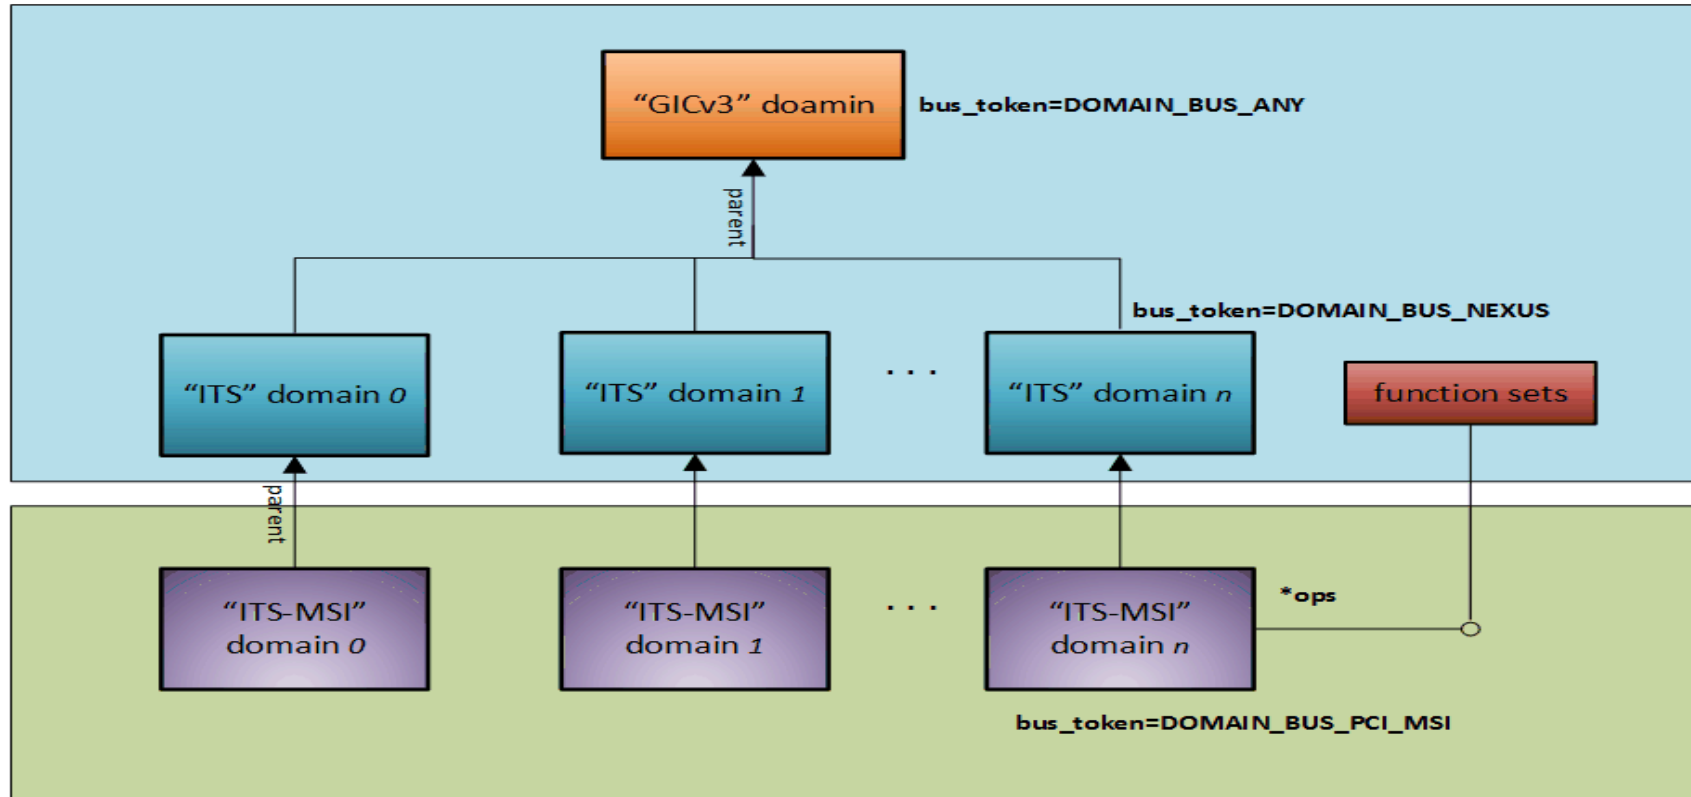
Field	Byte Length	Byte Offset	Description
Type	1	0	0xC GICD structure
Length	1	1	24
Reserved	2	2	Reserved - Must be zero
GIC ID	4	4	This GIC Distributor's hardware ID
Physical Base Address	8	8	The 64-bit physical address for this Distributor
System Vector Base	4	16	The global system interrupt number where this GIC Distributor's interrupt inputs start.
GIC version	1	20	0x00: No GIC version is specified, fall back to hardware discovery for GIC version 0x01: GICv1 0x02: GICv2 0x03: GICv3 0x04: GICv4 0x05-0xFF, Reserved for future use.
Reserved	3	21	Must be zero

`gic_acpi_init()` is used to parse the MADT and create the `irq_domain` hierarchy of GIC

# irq\_domain hierarchy of GIC

- irq\_domain level setup during GIC initialization

GIC logical domain

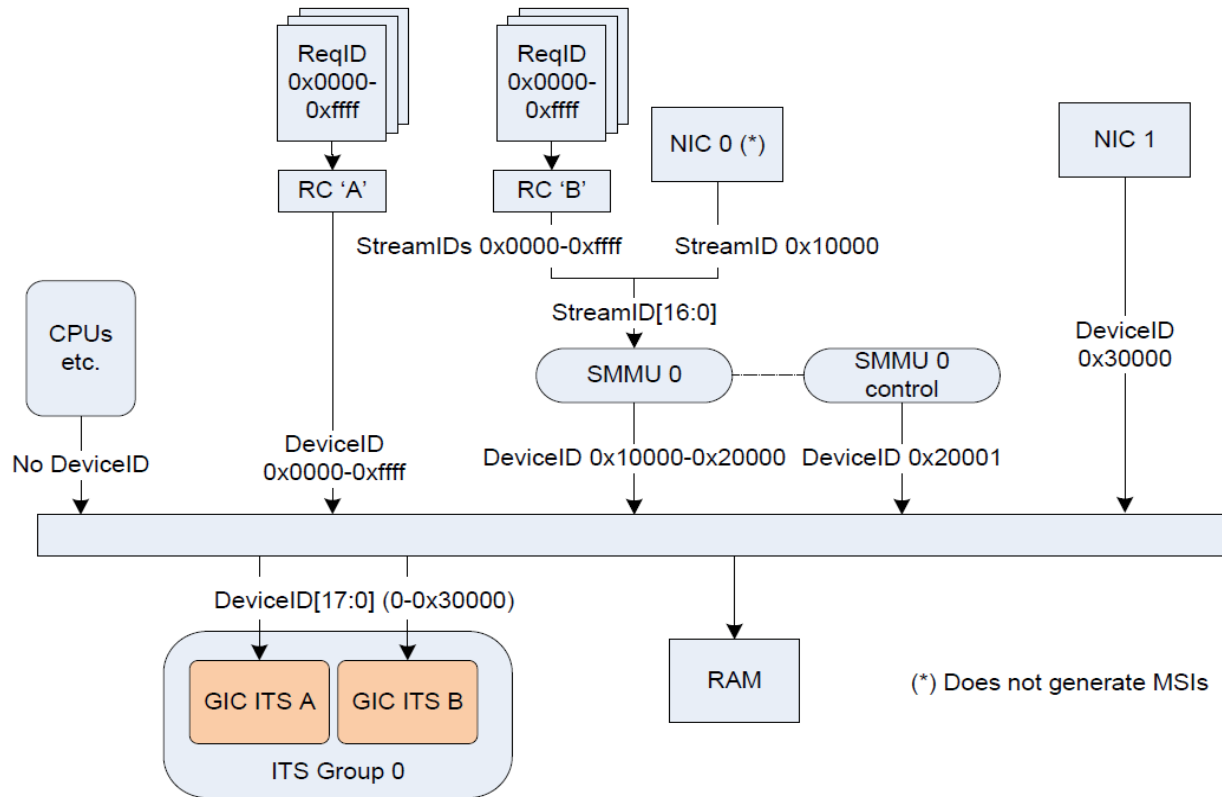


PCI MSI domain

GIC driver in kernel setups the above hierarchy by parsing the MADT



# DeviceID and IORT



Devices topology on AArch64 platform

- IORT is ACPI table to describe this kind of topology.
- GIC subsystem in kernel parses the IORT and generate the DeviceID.
- The mapping of the ReqID, StreamID and DeviceID is baked in Hardware, the IORT in firmware just describe this mapping, it can't modify the mapping relationship
- `iort_msi_map_rid()` used to map a ReqID to a DeviceID, then the GIC driver in kernel uses this DeviceID to setup a DTE for interrupt translation.

# Interrupts on AArch64

- Screenshot of interrupts on a 8-core AArch64 Fast Model

```
/ # cat /proc/interrupts
CPU0      CPU1      CPU2      CPU3      CPU4      CPU5      CPU6      CPU7
1:         0         0         0         0         0         0         0      GICv3  25 Level    vgic
2:         0         0         0         0         0         0         0      GICv3  29 Level    arch_timer
3:       569       292       320       206       447       221       198      GICv3  30 Level    arch_timer
4:         0         0         0         0         0         0         0      GICv3  27 Level    kvm guest timer
6:         0         0         0         0         0         0         0      GICv3  106 Edge    arm-smmu-v3-evtq
8:         0         0         0         0         0         0         0      GICv3  111 Edge    arm-smmu-v3-gerror
9:         0         0         0         0         0         0         0      GICv3  109 Edge    arm-smmu-v3-cmdq-sync
11:        10         0         0         0         0         0         0      GICv3  74 Level    virtio0
12:       389         0         0         0         0         0         0      GICv3  37 Level    uart-pl011
13:         0         0         0         0         0         0         0      ITS-MSI  0 Edge    virtio1-config
14:         0         0         0         0         0         0         0      ITS-MSI  1 Edge    virtio1-req.0
15:         0         0         0         0         0         0         0      ITS-MSI 32768 Edge    virtio2-config
16:         0         0         0         0         0         0         0      ITS-MSI 32769 Edge    virtio2-req.0
17:         2         0         0         0         0         0         0      ITS-MSI 49152 Edge    smmu-te
IPI0:      322      353      256      266      413      211      216      154      Rescheduling interrupts
IPI1:        5        6        5        6        4        1        5        6      Function call interrupts
IPI2:         0         0         0         0         0         0         0         0      CPU stop interrupts
IPI3:         0         0         0         0         0         0         0         0      Timer broadcast interrupts
IPI4:         0         0         0         0         0         0         0         0      IRQ work interrupts
IPI5:         0         0         0         0         0         0         0         0      CPU wake-up interrupts
Err:        0
/ #
```

*Take the smmu-te device as an example, irq =17 hwirq=49152, irq\_domain='ITS-MSI'.*

# Example of PCI MSI/MSI-X usage

- **struct msix\_entry**

```
struct msix_entry {  
    u32    vector;    // Filled by PCI MSI/MSIX subsystem, the irq  
    u16    entry;     // Filled by pci device driver, generally 0,1,... N. N is the total MSI-X vectorS numbers  
};
```

- **Enable the MSI/MSI-X**

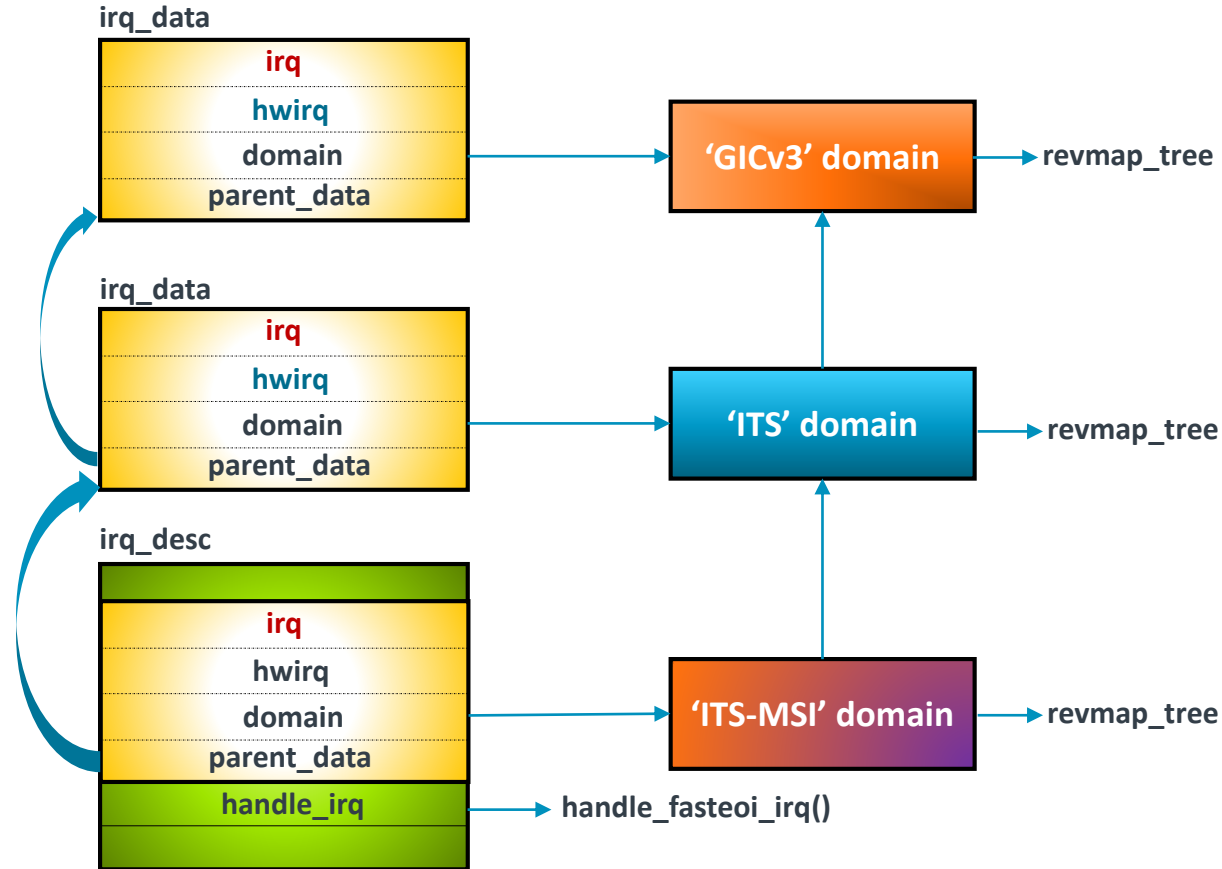
```
struct msix_entry *entries = devm_kmalloc(&pdev->dev, sizeof(struct msix_entry) * numvecs, GFP_KERNEL);  
for (i = 0; i < numvecs; i++)  
    entries[i].entry = i;  
...  
ret = pci_enable_msix_range(pdev, entries, 1, numvecs);  
  
if (ret <= 0)  
    return ret;
```

- **Register the interrupt handler**

```
for (i = 0; i < nr_rx; i++)  
    request_irq(entries[i].vector, rx_pkg_handler, 0, netdev->name, adapter);
```

# irq & hwirq Allocation

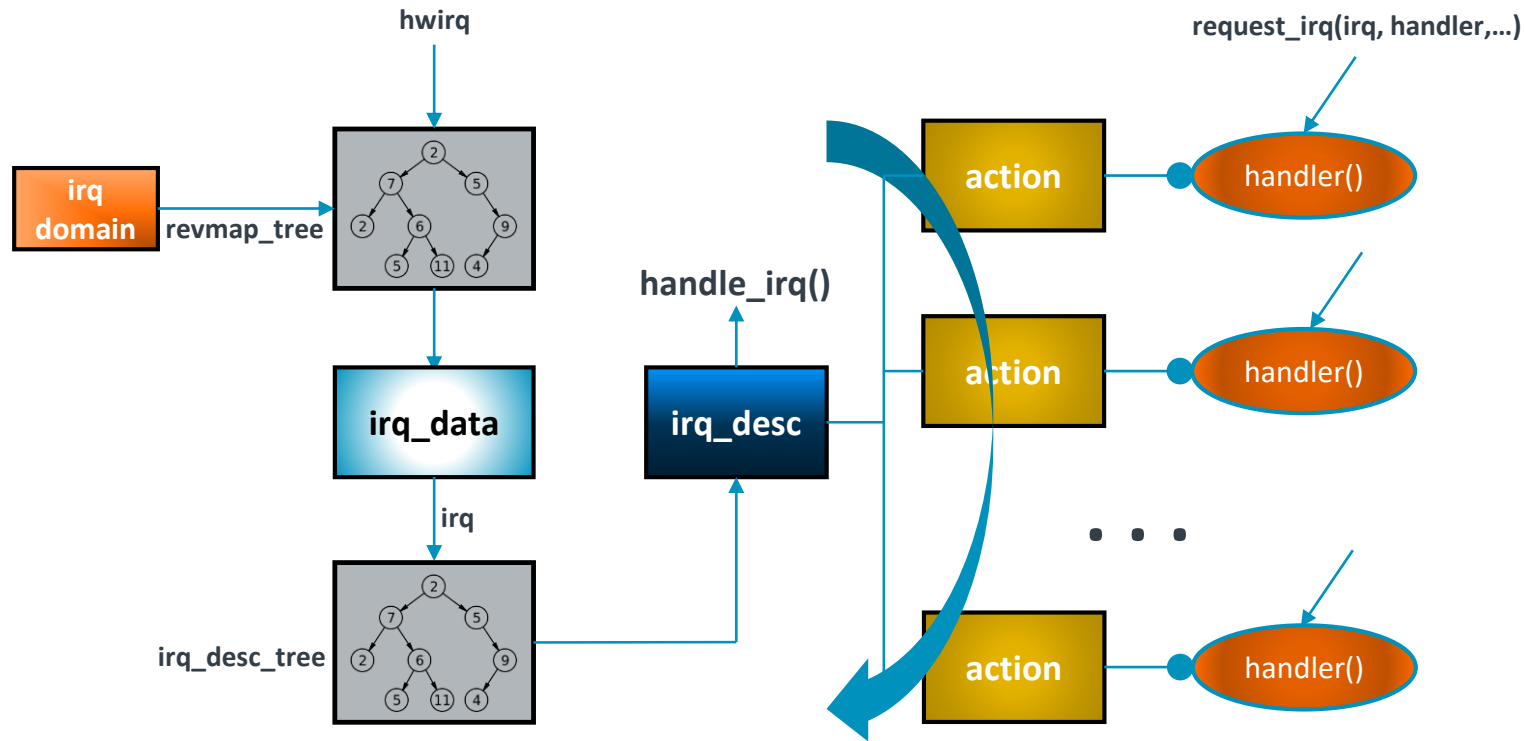
- Runtime allocation and mapping



- For a given pci device:  
`pci_msi_get_device_domain(pdev) →`  
`dom = iort_get_device_domain(&pdev->dev, rid)` will find the "ITS-MSI" domain according to the rid of the pdev
- The "ITS-MSI" domain is bound to a pci device(pdev) which happens in:  
`pci_device_add() → pci_set_msi_domain(pdev)`
- The irq is global in the irq\_domain space, hwirq is not. each domain will allocate a hwirq.
- each domain has a hwirq used as an index of the radix tree:  
`radix_tree_insert(&domain->revmap_tree, hwirq, irq_data)`

# Installation of interrupt handler

- How the interconnection setup



- Read the hwirq from the GIC (ICC\_IAR)
- Lookup the irq\_data from the revmap\_tree (indexed by hwirq)
- Lookup the irq\_desc from the irq\_desc\_tree (indexed by irq)
- Invoke irq\_desc->handle\_irq()
- Call the action->handler() iteratively

# Hardirq & Softirq

- **Hardirq part**

- Can't receive the external interrupt when in Hardirq – Behavior of the hardware  
PSTATE.{I, F} will be set 1 to mask external IRQ and FIQ
- Software need to enable the IRQ/FIQ at the end of Hardirq

- **Softirq part**

- Kernel unmask the external interrupt before invoking the softirq
- Core can receive external interrupt during the softirq handling
- Both hardirq and softirq are in the interrupt context

- **Arm instruction enable/disable IRQ**

```
static inline void arch_local_irq_enable(void)
{
    asm volatile("msr daifclr, #2"
                 ::: "memory");
}
```

```
static inline void arch_local_irq_disable(void)
{
    asm volatile("msr daifset, #2"
                 ::: "memory");
}
```

# PCI(e) MSI/MSI-X implementation

## ■ MSI-X Capability & Table structure




Figure 6-10: MSI-X Capability Structure




Figure 6-11: MSI-X Table Structure

(Table Base) bir = (u8)table\_offset & 0x7;

(Table Base) base = dev->resource[bir].start;


Table Phys Addr = base + table\_offset & FFFFFFFF8;

# PCI(e) MSI/MSI-X implementation

## ■ GIC ITS Translation Register

### Field descriptions

The GITS\_TRANSLATER bit assignments are:



**EventID, bits [31:0]**

An identifier corresponding to the interrupt to be translated.

pci\_enable\_msix\_range() → msi\_domain\_activate() →  
irq\_chip\_compose\_msi\_msg() → its\_irq\_compose\_msi\_msg()

```
static void its_irq_compose_msi_msg(struct irq_data *d, struct msi_msg *msg)
{
    struct its_device *its_dev = irq_data_get_irq_chip_data(d);
    struct its_node *its;
    u64 addr;

    its = its_dev->its;
    addr = its->phys_base + GITS_TRANSLATER;

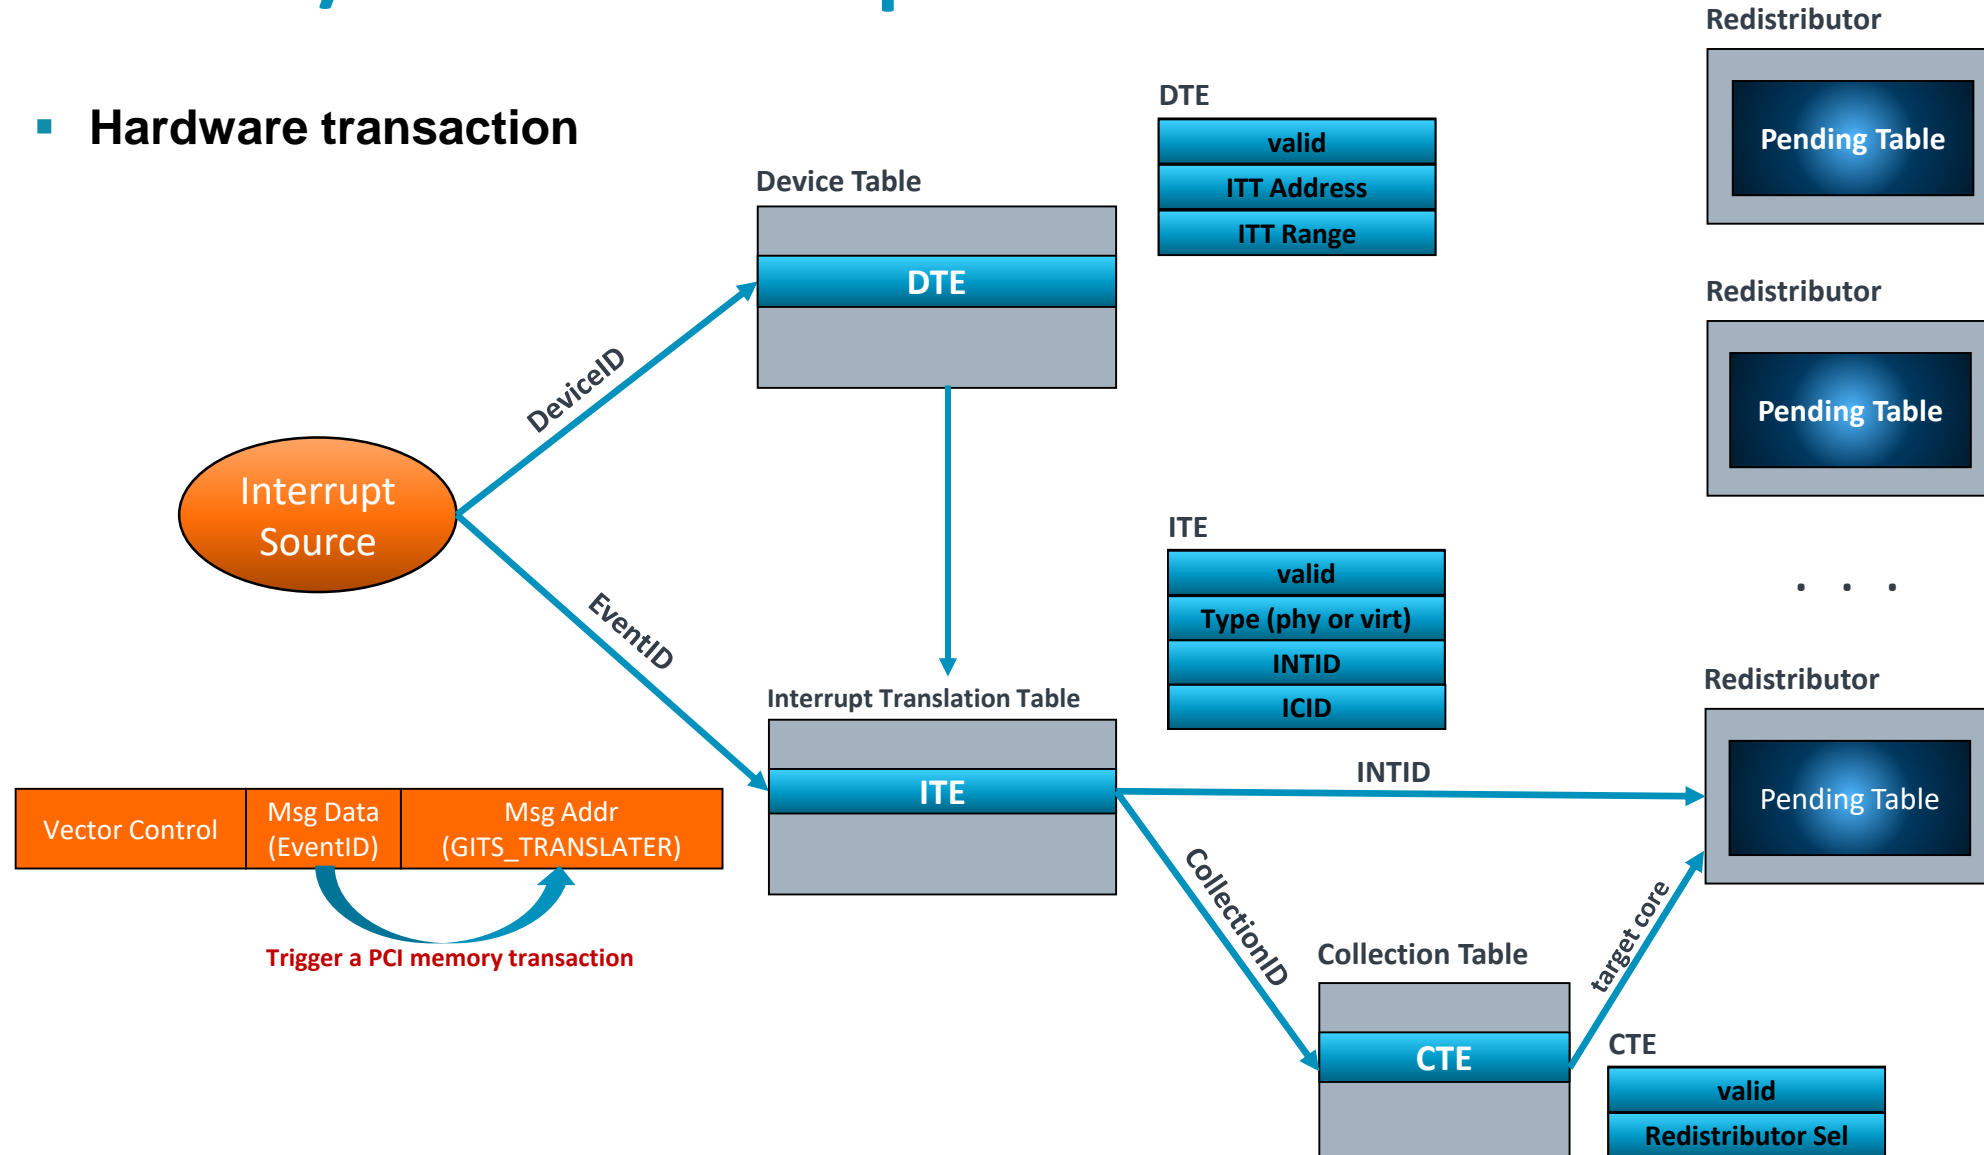
    msg->address_lo          = lower_32_bits(addr);
    msg->address_hi          = upper_32_bits(addr);
    msg->data                = its_get_event_id(d);

    iommu_dma_map_msi_msg(d->irq, msg);
}
```



# GIC MSI/MSI-X interrupt flow

- Hardware transaction



# Thank You





The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

[www.arm.com/company/policies/trademarks](http://www.arm.com/company/policies/trademarks)